

# Introduction to Database Systems

## CSE 414

### Lecture 4: SQL Joins and Aggregates

# Announcements

- Homework 2 out now
  - git pull upstream master to get the starter code
  - Due Tuesday Oct. 9 at midnight
- Web quiz 1 due Friday midnight
- Section tomorrow important for HW2

# Loading Data into SQLite

```
>sqlite3 lecture04.db  
sqlite> create table Purchase  
    (pid int primary key,  
     product text,  
     price float,  
     quantity int,  
     month varchar(15));  
sqlite> -- download data.txt  
sqlite> .mode list      always explicitly state, even if you think you are already in that mode (sanity check)  
sqlite> .import lec04-data.txt Purchase
```

Specify a filename where the database will be stored

Other DBMSs have other ways of importing data

## MUST BE IN SAME MODE AS FILE TYPE

If the data is separated by commas, need to set

```
sqlite> .mode csv
```

Product(pname, price, category, manufacturer)

Company(cname, country)

# Joins in SQL

pname	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all Japanese products that cost < \$150

```
SELECT P.pname, P.price  
FROM Product as P, Company as C  
WHERE P.manufacturer=C.cname AND  
C.country='Japan' AND C.price < 150
```

Join Predicate

Selection predicates

Product(pname, price, category, manufacturer)

Company(cname, country)

# Joins in SQL

pname	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all USA companies  
that manufacture “gadget” products

Product(pname, price, category, manufacturer)

Company(cname, country)

# Joins in SQL

pname	price	category	manufacturer
MultiTouch	199.99	gadget	Canon
SingleTouch	49.99	photography	Canon
Gizom	50	gadget	GizmoWorks
SuperGizmo	250.00	gadget	GizmoWorks

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

Retrieve all USA companies  
that manufacture “gadget” products

```
SELECT DISTINCT C.cname
FROM Product as P, Company as C
WHERE C.country='USA' AND P.category = 'gadget'
AND P.manufacturer = cname
```

Why  
**DISTINCT?**

# (Inner) joins

```
Product(pname, price, category, manufacturer)
Company(cname, country)
-- manufacturer is foreign key to Company
```

```
SELECT DISTINCT cname
FROM Product, Company
WHERE country='USA' AND category = 'gadget'
      AND manufacturer = cname
```

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
AND manufacturer = cname
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
Camera	Photo	Hitachi
OneClick	Photo	Hitachi

Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
AND manufacturer = cname
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
Camera	Photo	Hitachi
OneClick	Photo	Hitachi

Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
AND manufacturer = cname
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
Camera	Photo	Hitachi
OneClick	Photo	Hitachi

Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
AND manufacturer = cname
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
Camera	Photo	Hitachi
OneClick	Photo	Hitachi

Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

pname	category	manufacturer	cname	country
Gizmo	gadget	GizmoWorks	GizmoWorks	USA

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
AND manufacturer = cname
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
Camera	Photo	Hitachi
OneClick	Photo	Hitachi

Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
AND manufacturer = cname
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
Camera	Photo	Hitachi
OneClick	Photo	Hitachi

Company

cname	country
GizmoWorks	USA
Canon	Japan
Hitachi	Japan

# (Inner) joins

```
SELECT DISTINCT cname  
FROM Product, Company  
WHERE country='USA' AND category = 'gadget'  
      AND manufacturer = cname
```

```
SELECT DISTINCT cname  
FROM Product JOIN Company ON  
country = 'USA' AND category = 'gadget'  
      AND manufacturer = cname
```

# (Inner) Joins

```
SELECT  x1.a1, x2.a2, ... xm.am  
FROM    R1 as x1, R2 as x2, ... Rm as xm  
WHERE   Cond
```

for  $x_1$  in  $R_1$ :

  for  $x_2$  in  $R_2$ :

  ...

    for  $x_m$  in  $R_m$ :

      if Cond( $x_1, x_2 \dots$ ):

        output( $x_1.a_1, x_2.a_2, \dots x_m.a_m$ )

This is called nested loop semantics since we are interpreting what a join means using a nested loop

# Another example

Product(pname, price, category, manufacturer)  
Company(cname, country)

-- manufacturer is foreign key to Company

Find US companies that manufacture both  
'gadgets' and 'photo' products

# Another example

Product(pname, price, category, manufacturer)  
Company(cname, country)

-- manufacturer is foreign key to Company

Find US companies that manufacture both  
'gadgets' and 'photo' products

```
SELECT DISTINCT z.cname
FROM Product x, Company z
WHERE z.country = 'USA'
    AND x.manufacturer = z.cname
    AND x.category = 'gadget'
    AND x.category = 'photography';
```

Does this work?

no because no single tuples will have both gadget and photography in category

# Another example

Product(pname, price, category, manufacturer)  
Company(cname, country)

-- manufacturer is foreign key to Company

Find US companies that manufacture **both**  
'gadgets' and 'photo' products

```
SELECT DISTINCT z.cname
FROM Product x, Company z
WHERE z.country = 'USA'
    AND x.manufacturer = z.cname
    AND (x.category = 'gadget'
        OR x.category = 'photography');
```

What about  
this?

no because this only checks if the tuple has either photo or gadget, not both.

# Another example

```
Product(pname, price, category, manufacturer)  
Company(cname, country)
```

-- manufacturer is foreign key to Company

Find US companies that manufacture both  
'gadgets' and 'photo' products

```
SELECT DISTINCT z.cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
AND x.manufacturer = z.cname  
AND y.manufacturer = z.cname  
AND x.category = 'gadget'  
AND y.category = 'photography';
```

Need to include  
Product twice!

# Self-Joins and Tuple Variables

Find US companies that manufacture both  
‘gadgets’ and ‘photo’ products

- Joining Product with Company is insufficient: need to join Product, with Product, and with Company
- When a relation occurs twice in the FROM clause we call it a self-join; in that case we must use tuple variables (aka table aliases) (why?)

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

pname	category	manufacturer
Gizmo	gadget	GizmoWorks
SingleTouch	photo	Hitachi
MultiTouch	Photo	GizmoWorks

Company

cname	country
GizmoWorks	USA
Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
	Gizmo	gadget	GizmoWorks
	SingleTouch	photo	Hitachi
	MultiTouch	Photo	GizmoWorks

Company

cname	country
GizmoWorks	USA
Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
y	Gizmo	gadget	GizmoWorks
SingleTouch	photo	Hitachi	
MultiTouch	Photo	GizmoWorks	

Company

cname	country
GizmoWorks	USA
Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

x	pname	category	manufacturer
y	Gizmo	gadget	GizmoWorks
	SingleTouch	photo	Hitachi
	MultiTouch	Photo	GizmoWorks

Company

z	cname	country
	GizmoWorks	USA
	Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
y	Gizmo	gadget	GizmoWorks
SingleTouch	photo	Hitachi	
MultiTouch	Photo	GizmoWorks	

Company Z

cname	country
GizmoWorks	USA
Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
	Gizmo	gadget	GizmoWorks
y	SingleTouch	photo	Hitachi
	MultiTouch	Photo	GizmoWorks

Company

Z	cname	country
	GizmoWorks	USA
	Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
	Gizmo	gadget	GizmoWorks
y	SingleTouch	photo	Hitachi
	MultiTouch	Photo	GizmoWorks

Company

Z	cname	country
	GizmoWorks	USA
	Hitachi	Japan

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
    AND x.category = 'gadget'  
    AND y.category = 'photo'  
    AND x.manufacturer = z cname  
    AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
	Gizmo	gadget	GizmoWorks
	SingleTouch	photo	Hitachi
y	MultiTouch	Photo	GizmoWorks

Company

Z	cname	country
	GizmoWorks	USA
	Hitachi	Japan

# Self-joins

```

SELECT DISTINCT z cname
FROM Product x, Product y, Company z
WHERE z.country = 'USA'
    AND x.category = 'gadget'
    AND y.category = 'photo'
    AND x.manufacturer = z cname
    AND y.manufacturer = z cname;

```

Product

X	pname	category	manufacturer
	Gizmo	gadget	GizmoWorks
	SingleTouch	photo	Hitachi
y	MultiTouch	Photo	GizmoWorks

Company

Z	cname	country
	GizmoWorks	USA
	Hitachi	Japan

x.pname	x.category	x.manufacturer	y.pname	y.category	y.manufacturer	z cname	z.country
Gizmo	gadget	GizmoWorks	MultiTouch	Photo	GizmoWorks	GizmoWorks	29 USA

# Self-joins

```
SELECT DISTINCT z cname  
FROM Product x, Product y, Company z  
WHERE z.country = 'USA'  
AND x.category = 'gadget'  
AND y.category = 'photo'  
AND x.manufacturer = z cname  
AND y.manufacturer = z cname;
```

Product

X	pname	category	manufacturer
	Gizmo	gadget	GizmoWorks
	SingleTouch	photo	Hitachi
y	MultiTouch	Photo	GizmoWorks

Company

Z	cname	country
	GizmoWorks	USA
	Hitachi	Japan

x.pname	x.category	x.manufacturer	y.pname	y.category	y.manufacturer	z cname	z.country
Gizmo	gadget	GizmoWorks	MultiTouch	Photo	GizmoWorks	GizmoWorks	30 USA

# Joins in SQL

- The join we have just seen is sometimes called an **inner join**
  - Each row in the result **must come from both tables in the join**
- Sometimes we want to include rows from only one of the two table: **outer join**

`Employee(id, name)`

`Sales(employeeID, productID)`

# Inner Join

Employee

<u>id</u>	<u>name</u>
1	Joe
2	Jack
3	Jill

Sales

<u>employeeID</u>	<u>productID</u>
1	344
1	355
2	544

Retrieve employees and their sales

`Employee(id, name)`

`Sales(employeeID, productID)`

# Inner Join

Employee

<u>id</u>	<u>name</u>
1	Joe
2	Jack
3	Jill

Sales

<u>employeeID</u>	<u>productID</u>
1	344
1	355
2	544

Retrieve employees and their sales

```
SELECT *
FROM   Employee E, Sales S
WHERE  E.id = S.employeeID
```

`Employee(id, name)`  
`Sales(employeeID, productID)`

# Inner Join

Employee

<u>id</u>	<u>name</u>
1	Joe
2	Jack
3	Jill

Sales

<u>employeeID</u>	<u>productID</u>
1	344
1	355
2	544

Retrieve employees and their sales

```
SELECT *
FROM Employee E, Sales S
WHERE E.id = S.employeeID
```

<u>id</u>	<u>name</u>	<u>employeeID</u>	<u>productID</u>
1	Joe	1	344
1	Joe	1	355
2	Jack	2	544

`Employee(id, name)`  
`Sales(employeeID, productID)`

# Inner Join

Employee

<u>id</u>	<u>name</u>
1	Joe
2	Jack
3	Jill

Sales

<u>employeeID</u>	<u>productID</u>
1	344
1	355
2	544

Retrieve employees and their sales

```
SELECT *
FROM Employee E, Sales S
WHERE E.id = S.employeeID
```



<u>id</u>	<u>name</u>	<u>employeeID</u>	<u>productID</u>
1	Joe	1	344
1	Joe	1	355
2	Jack	2	544

```
Employee(id, name)  
Sales(employeeID, productID)
```

# Inner Join

Employee

id	name
1	Joe
2	Jack
3	Jill

Sales

employeeID	productID
1	344
1	355
2	544

Retrieve employees and their sales

Alternative syntax

```
SELECT *  
FROM Employee E  
INNER JOIN  
Sales S  
ON E.id = S.employeeID
```

Jill is missing

id	name	employeeID	productID
1	Joe	1	344
1	Joe	1	355
2	Jack	2	544

`Employee(id, name)`  
`Sales(employeeID, productID)`

# Outer Join

can just think of the same nested loop semantics as with inner join, but add default behavior where if a row/tuple in first/second table (based on Left or right outer join respectively) has no matches, simply fill in the missing entries with NULLs

Employee

<u>id</u>	<u>name</u>
1	Joe
2	Jack
3	Jill

Sales

<u>employeeID</u>	<u>productID</u>
1	344
1	355
2	544

Retrieve employees and their sales

Jill is present

```
SELECT *
FROM Employee E
LEFT OUTER JOIN
Sales S
ON E.id = S.employeeID
```

<u>id</u>	<u>name</u>	<u>employeeID</u>	<u>productID</u>
1	Joe	1	344
1	Joe	1	355
2	Jack	2	544
3	Jill	NULL	NULL

# Outer joins

Product(name, category)

Purchase(prodName, store)

-- prodName is foreign key

```
SELECT Product.name, Purchase.store  
FROM Product LEFT OUTER JOIN Purchase ON  
Product.name = Purchase.prodName
```

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz

```

SELECT Product.name, Purchase.store
FROM Product JOIN Purchase ON
Product.name = Purchase.prodName

```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz
Camera	Ritz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```
SELECT Product.name, Purchase.store  
FROM Product JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```

SELECT Product.name, Purchase.store
FROM Product LEFT OUTER JOIN Purchase ON
Product.name = Purchase.prodName

```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

```

SELECT Product.name, Purchase.store
FROM Product LEFT OUTER JOIN Purchase ON
Product.name = Purchase.prodName

```

could use  
WHERE  
here instead

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Output

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL

FULL outer join does both a left/right join. i.e. any tuples in either table which have no match will be added with NULL entries

```
SELECT Product.name, Purchase.store  
FROM Product FULL OUTER JOIN Purchase ON  
Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
Phone	Foo

Output

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL
NULL	Foo

# Outer Joins

```
tableA (LEFT/RIGHT/FULL) OUTER JOIN tableB ON p
```

- Left outer join:
  - Include tuples from tableA even if no match
- Right outer join:
  - Include tuples from tableB even if no match
- Full outer join:
  - Include tuples from both even if no match
- In all cases:
  - Patch tuples without matches using NULL

# Aggregates in SQL

# Simple Aggregations

Five basic aggregate operations in SQL

```
select count(*) from Purchase counts the number of tuples that you have selected
select sum(quantity) from Purchase sums values of “quantity” attribute for selected entries
select avg(price) from Purchase
select max(quantity) from Purchase
select min(quantity) from Purchase
```

Except count, all aggregations apply to a **single attribute**

# Aggregates and NULL Values

Null values are not used in aggregates

```
insert into Purchase  
values(12, 'gadget', NULL, NULL, 'april')
```

Try the following

```
select count(*) from Purchase
```

counts the number of tuples, regardless  
of whether they have null entries or not

```
select count(quantity) from Purchase
```

would count tuple as long as  
quantity attribute is not Null

```
select sum(quantity) from Purchase
```

```
select count(*)
```

```
from Purchase
```

```
where quantity is not null;
```

# Counting Duplicates

COUNT applies to duplicates, unless otherwise stated:

```
SELECT count(product)
FROM Purchase
WHERE price > 4.99
```

same as count(\*) if no nulls

We probably want:

```
SELECT count(DISTINCT product)
FROM Purchase
WHERE price > 4.99
```

# More Examples

```
SELECT Sum(P.price * P.quantity)  
FROM Purchase as P
```

What do  
they mean ?

```
SELECT Sum(P.price * P.quantity)  
FROM Purchase as P  
WHERE P.product = 'bagel'
```