

Introduction to Database Systems

CSE 414

Lecture 11: NoSQL

Announcements

- HW 3 due Friday
 - Upload data with DataGrip editor – see message board
 - Azure timeout for question 5:
 - Try DataGrip or SQLite
- HW 2 Grades and Feedback out
 - Check feedback, some tag errors
- HW 4 posted today, due week from Tuesday

Class Overview

- Unit 1: Intro
- Unit 2: Relational Data Models and Query Languages
- Unit 3: Non-relational data
 - NoSQL
 - Json
 - SQL++
- Unit 4: RDMBS internals and query optimization
- Unit 5: Parallel query processing
- Unit 6: DBMS usability, conceptual design
- Unit 7: Transactions
- Unit 8: Advanced topics (time permitting)

Two Classes of Database Applications

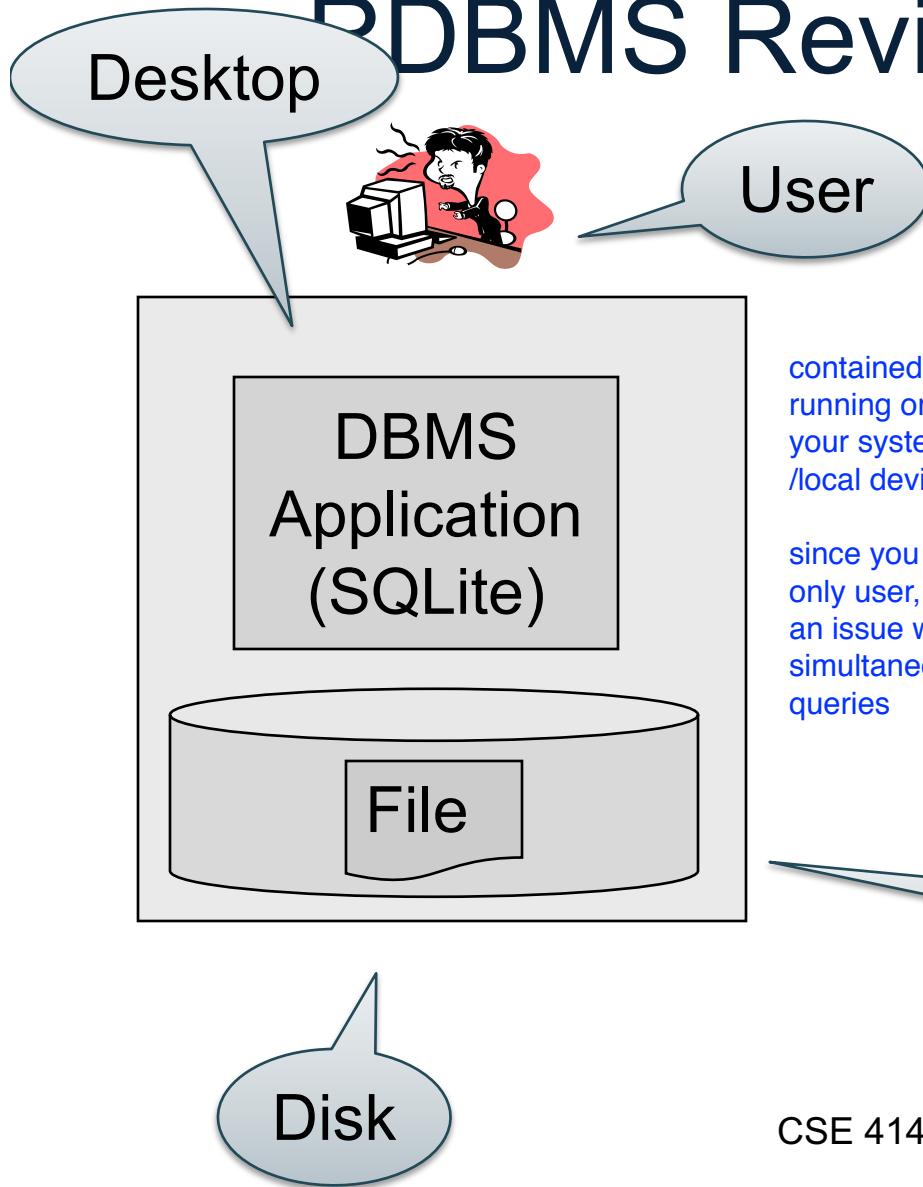
- OLTP (Online Transaction Processing)
 - Queries are simple lookups: 0 or 1 join
E.g., find customer by ID and their orders
 - Many updates. E.g., insert order, update payment
 - **Consistency** is critical: **transactions** (more later)
- OLAP (Online Analytical Processing)
 - aka “Decision Support”
 - Queries have many joins, and group-by’s
E.g., sum revenues by store, product, clerk, date
 - No updates

can we maintain a valid state if two identical queries come simultaneously?

NoSQL Motivation

- Originally motivated by Web 2.0 applications
 - E.g. Facebook, Amazon, Instagram, etc
 - Web startups need to scaleup from 10 to 100000 users very quickly
- Needed: very large scale OLTP workloads
- Give up on consistency
- Give up OLAP

RDBMS Review: Serverless



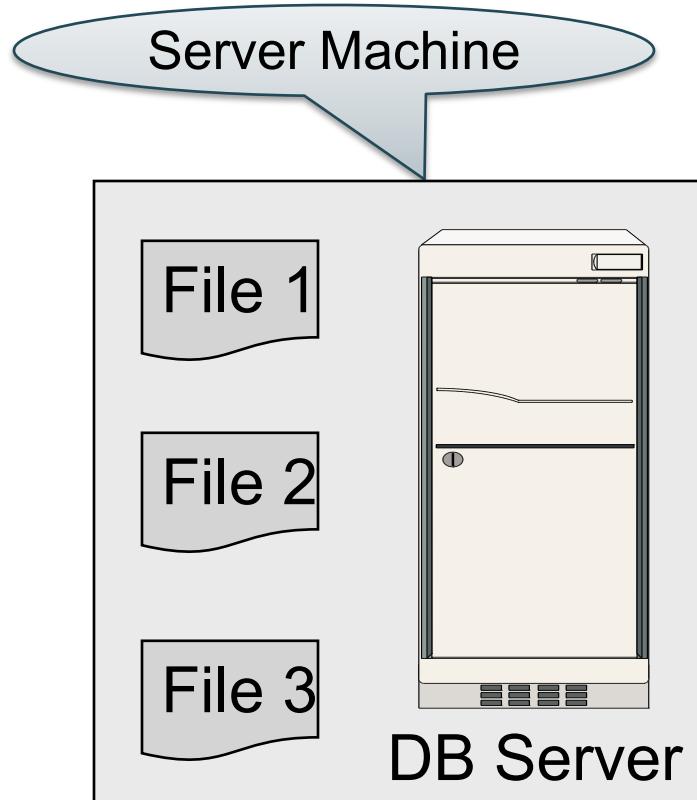
SQLite:

- One data file
- One user
- One DBMS application

• **Consistency** is easy

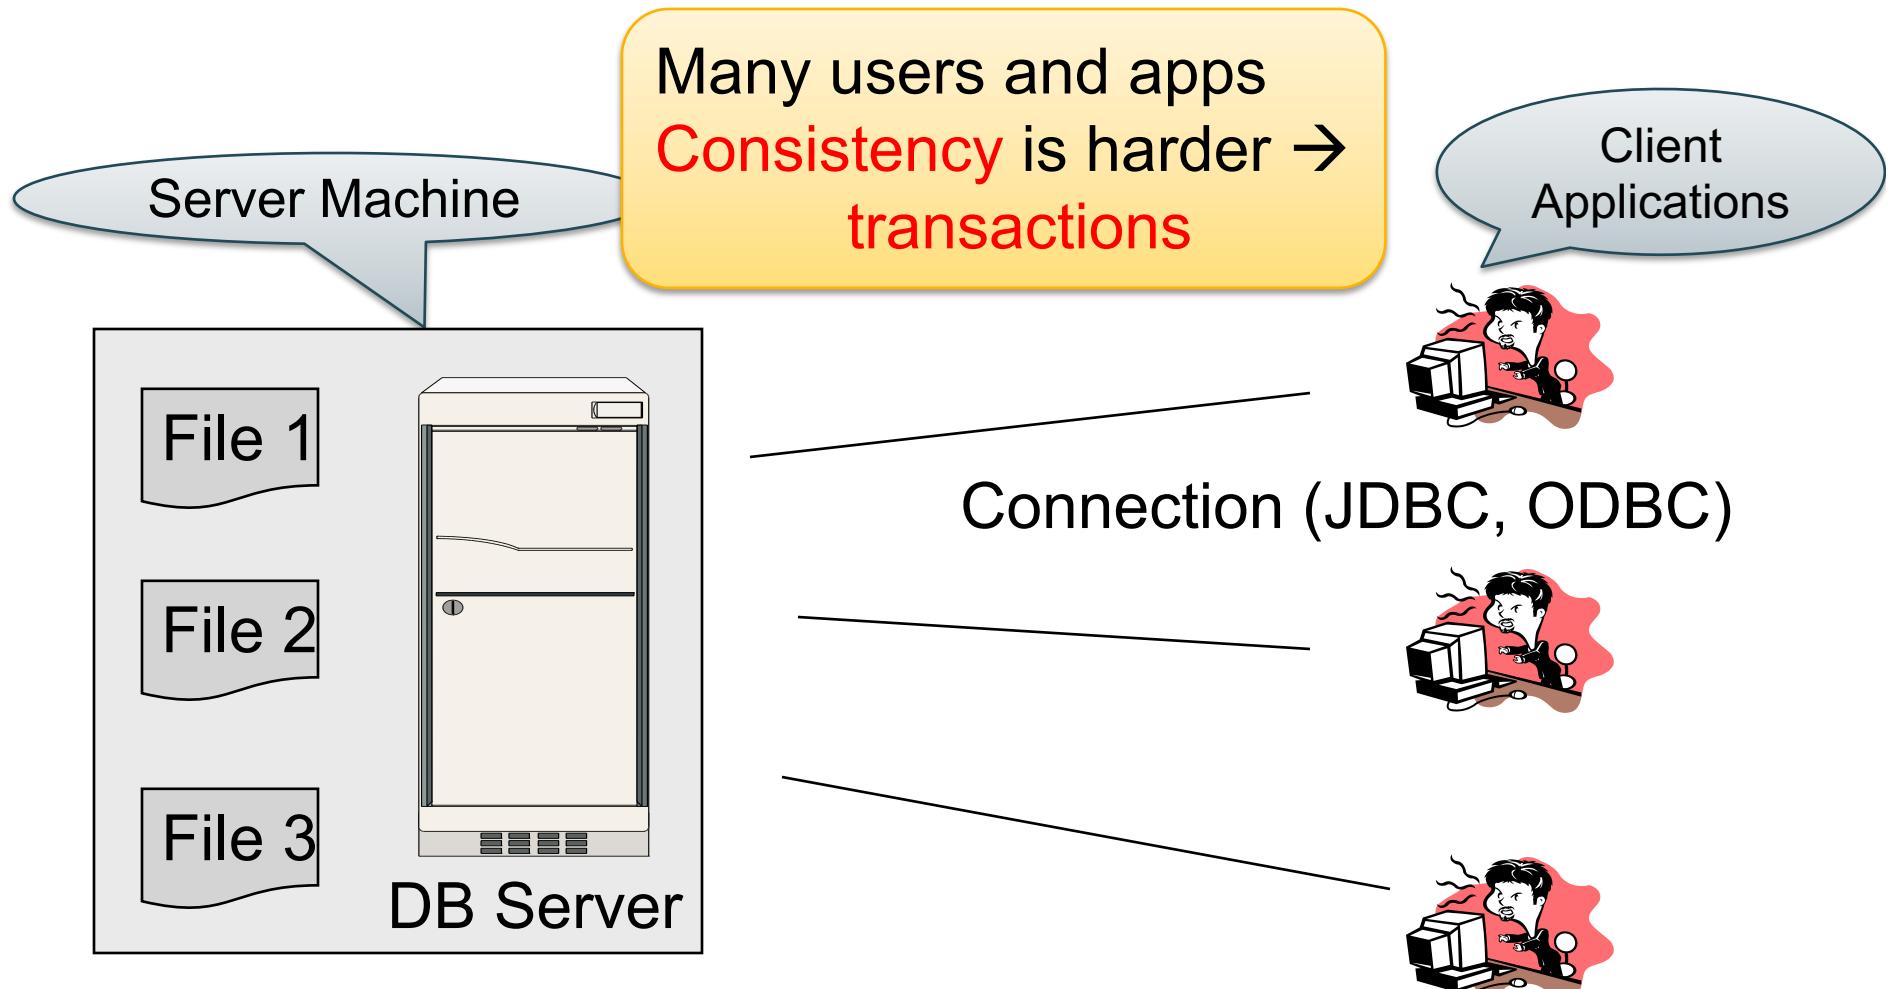
But only a limited number of scenarios work with such model

RDBMS Review: Client-Server



- One server running the database
- Many clients, connecting via the ODBC or JDBC (Java Database Connectivity) protocol

RDBMS Review: Client-Server



- One server running the database
- Many clients, connecting via the ODBC or JDBC (Java Database Connectivity) protocol

8

Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)

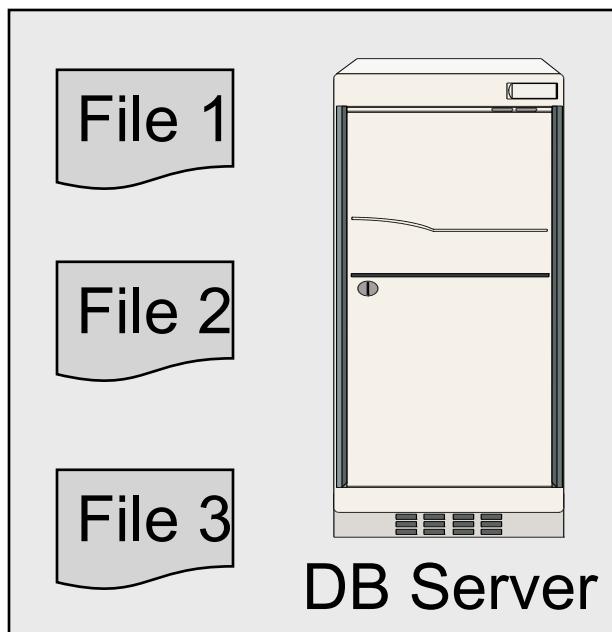
Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)
- Many *clients* run apps and connect to DBMS
 - Microsoft's Management Studio (for SQL Server), or
 - psql (for postgres)
 - Some Java program (HW8) or some C++ program

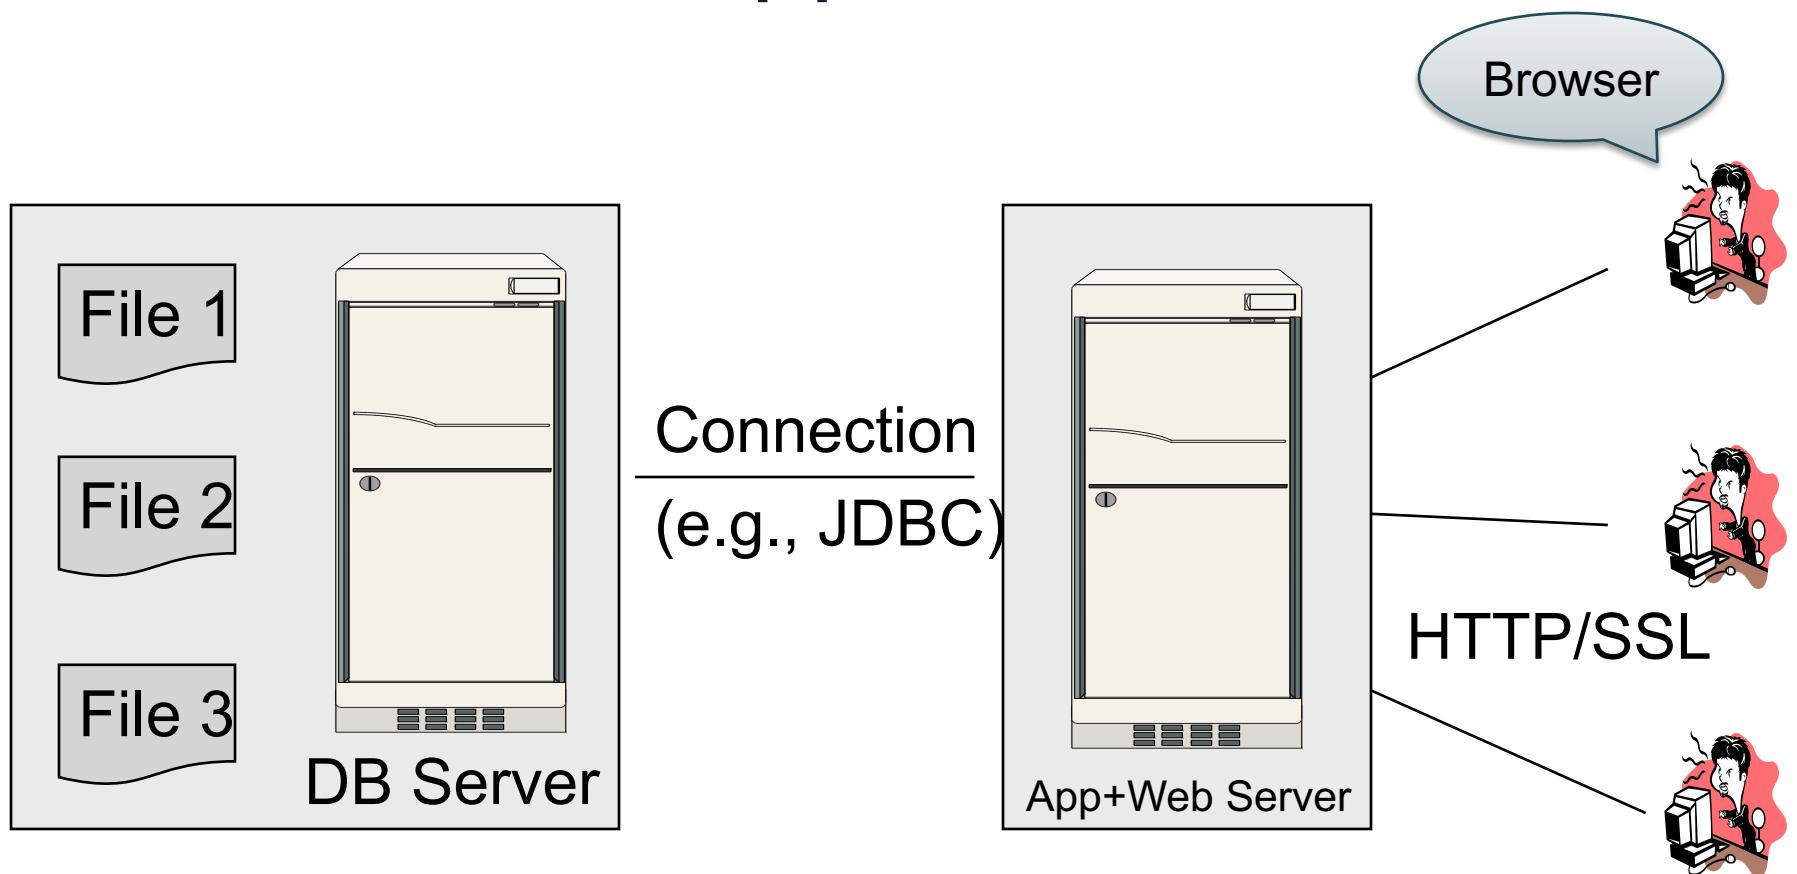
Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)
- Many *clients* run apps and connect to DBMS
 - Microsoft's Management Studio (for SQL Server), or
 - psql (for postgres)
 - Some Java program (HW8) or some C++ program
- Clients “talk” to server using JDBC/ODBC protocol

Web Apps: 3 Tier

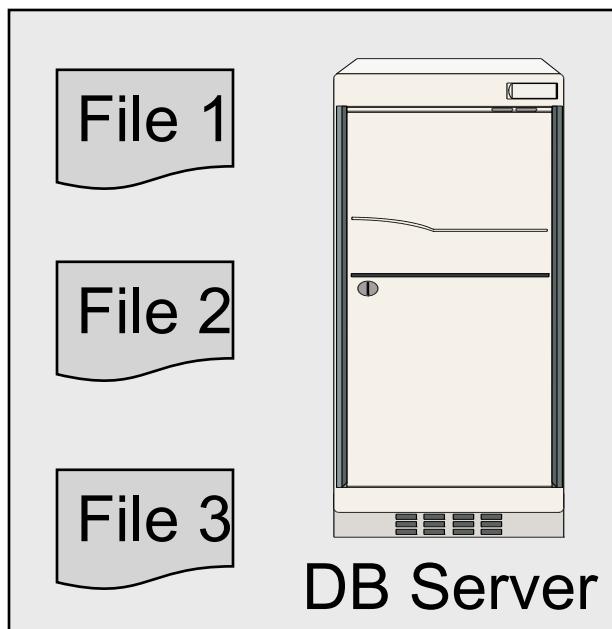


Web Apps: 3 Tier

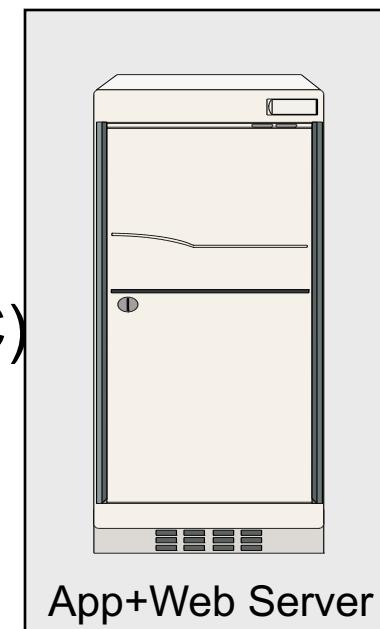


Web Apps: 3 Tier

Web-based applications



Connection
(e.g., JDBC)



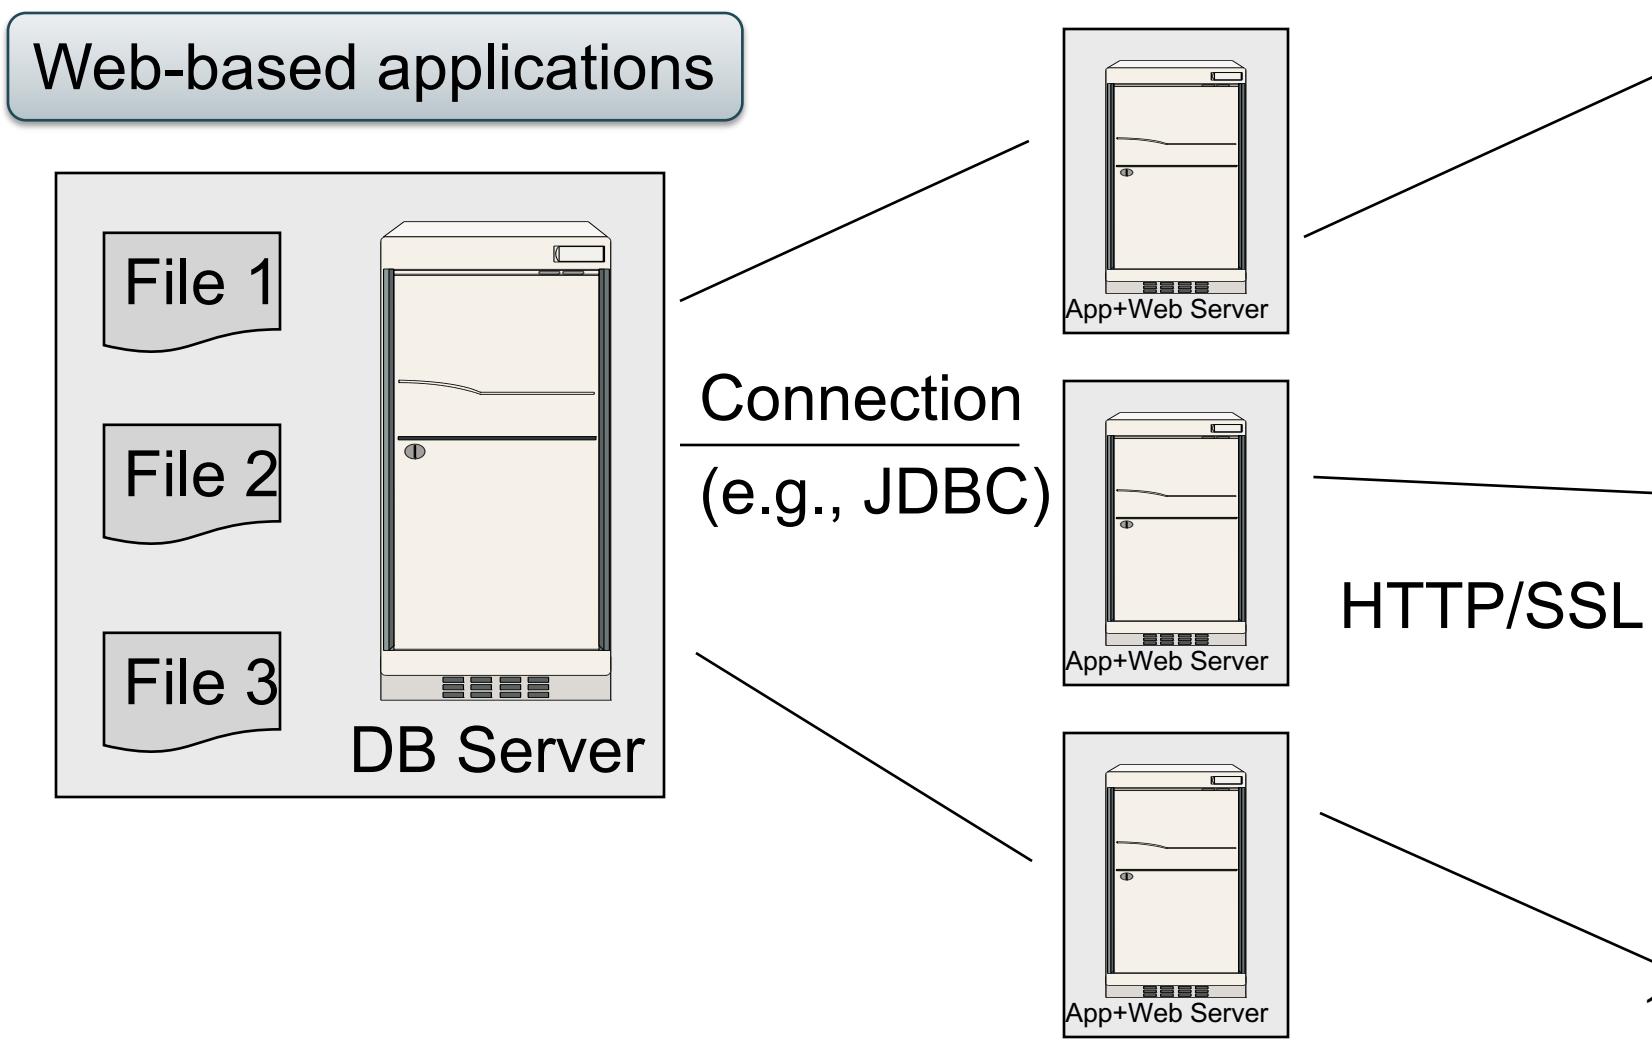
Browser



HTTP/SSL

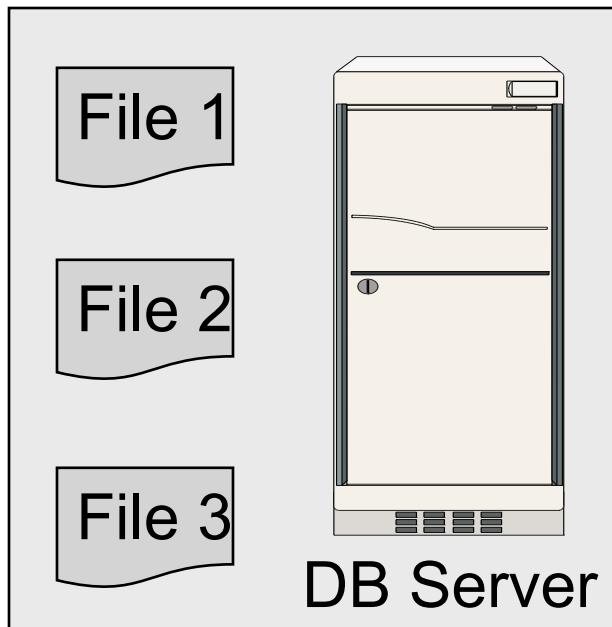


Web Apps: 3 Tier

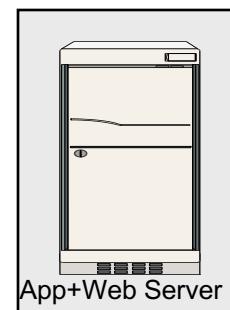
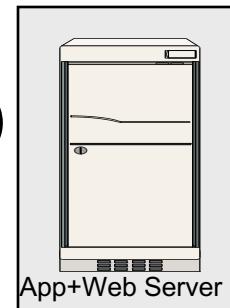
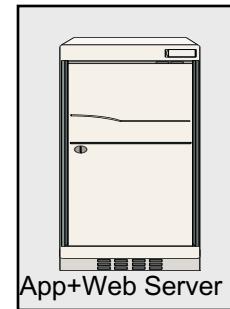


Replicate
App server
for scaleup

Web-based applications



Connection
(e.g., JDBC)

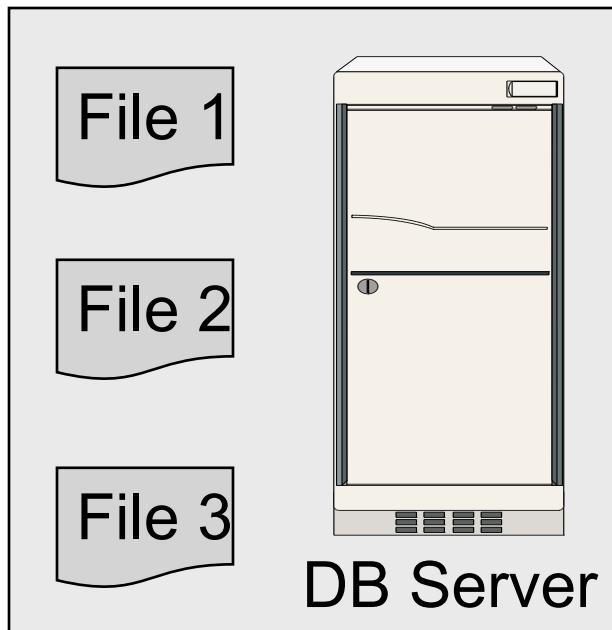


HTTP/SSL

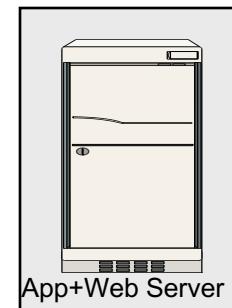
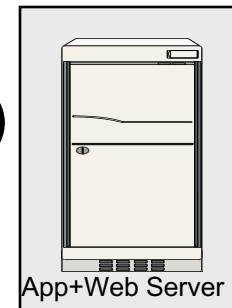
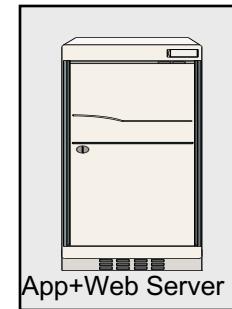
Why not replicate DB server?

Replicate
App server
for scaleup

Web-based applications



OS: 3 Tier



HTTP/SSL

Why not replicate DB server?
Consistency!

hard to keep files identical across both devices



Replicating the Database

INCREASE scale of the database (how many users it can handle)

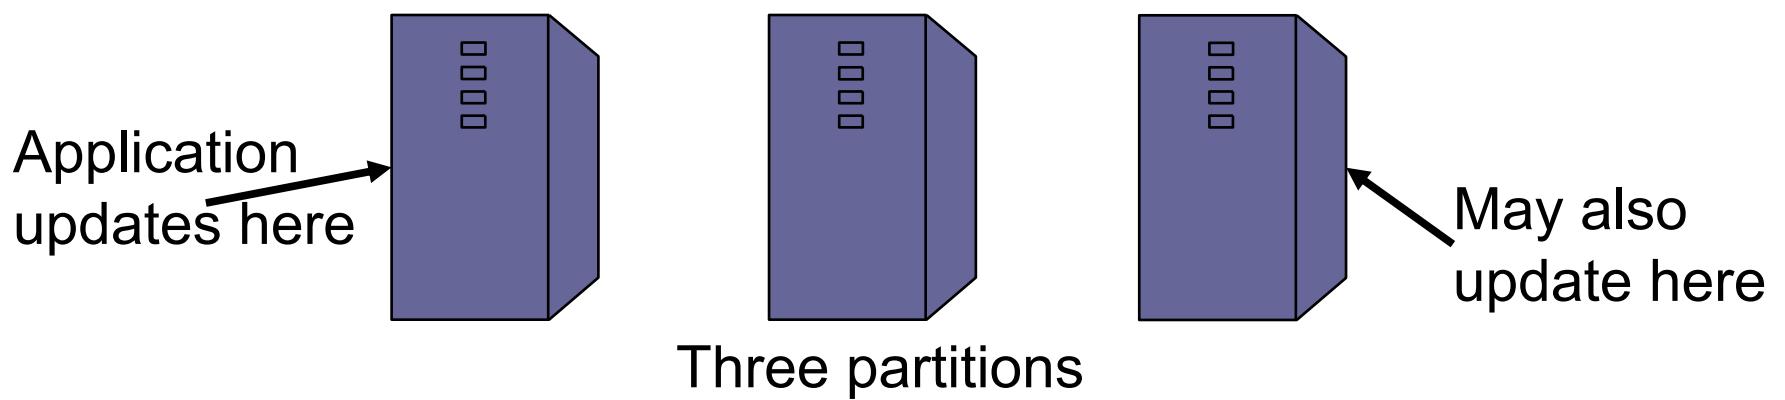
- Two basic approaches:
 - Scale up through **partitioning**
divide up data across multiple machines
 - Scale up through **replication**
multiple identical copies of database
- **Consistency** is much harder to enforce

Scale Through Partitioning

- Partition the database across many machines in a cluster
 - Database now fits in main memory
 - Queries spread across these machines
- Can increase throughput
- Easy for writes but reads become expensive!

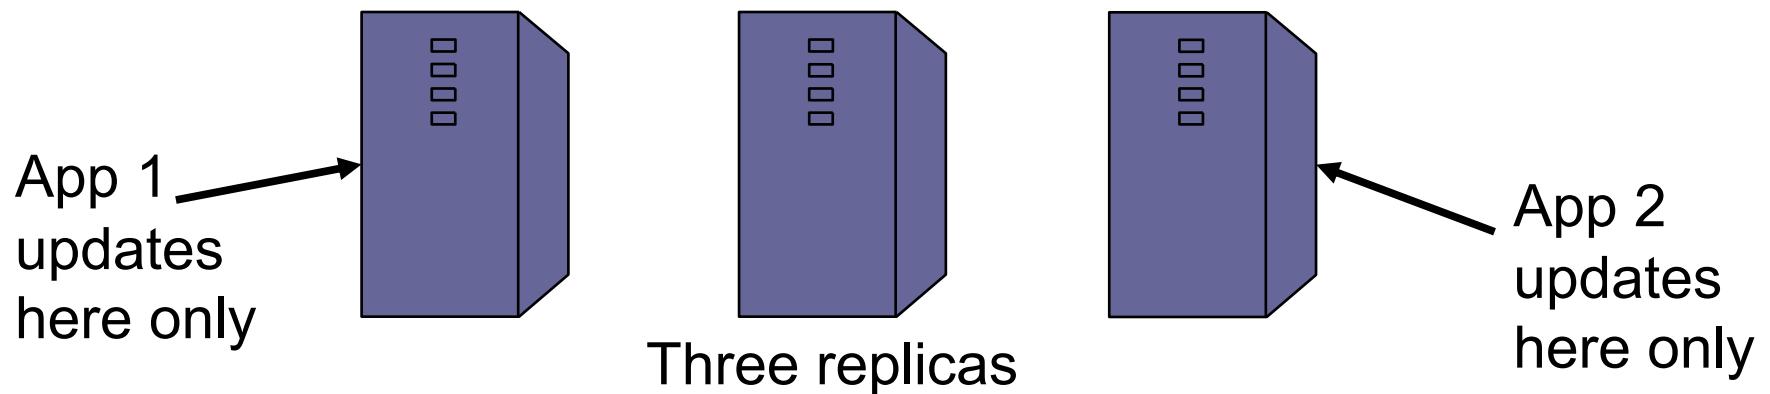
store individual attributes on each server?
Lowers space requirement per tuple, can store
whole database on a single machine

might have to read from every machine!



Scale Through Replication

- Create multiple copies of each database partition
- Spread queries across these replicas
- Can increase throughput and lower latency
- Can also improve fault-tolerance
as that piece of data is stored on separate machines; might have to talk to all machines
- Easy for reads but writes become expensive!



Relational Model → NoSQL

that's why a new data model was needed

- Relational DB: difficult to replicate/partition
- Given Supplier(sno,...), Part(pno,...), Supply(sno,pno)
 - Partition: we may be forced to join across servers
 - Replication: local copy has inconsistent versions
 - Consistency is hard in both cases (why?)
- NoSQL: simplified data model
 - Given up on functionality
 - Application must now handle joins and consistency

can only extract full tuples of data

Data Models

Taxonomy based on data models:

- ☞ • **Key-value stores** similar to hash tables
 - e.g., Project Voldemort, Memcached
- **Document stores** key value basically where value is a document
 - e.g., SimpleDB, CouchDB, MongoDB

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - `get(key)`, `put(key,value)`
 - Operations on value not supported

Key-Value Stores Features

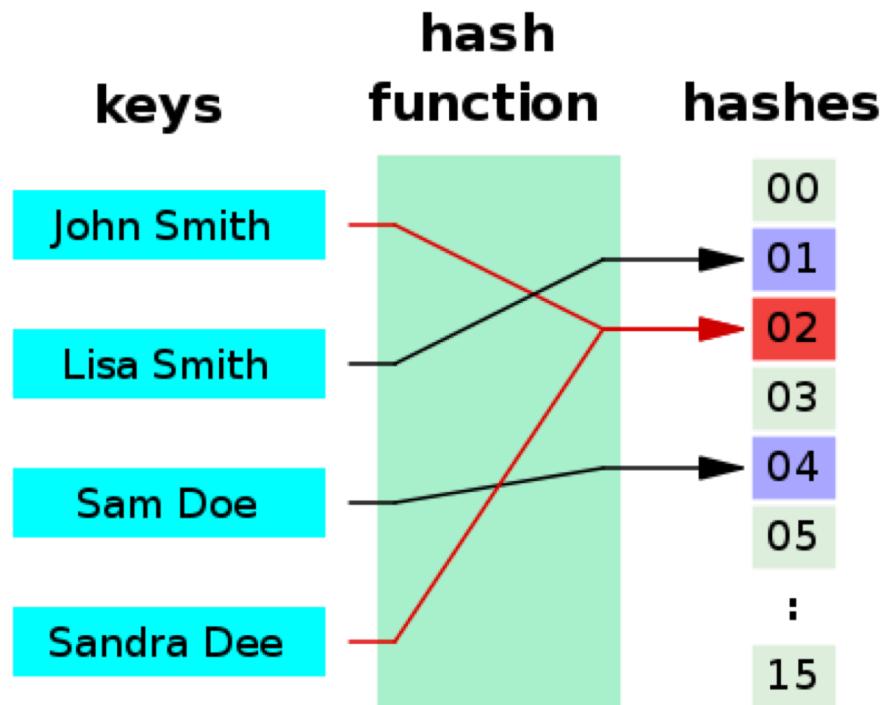
- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - get(key), put(key,value)
 - Operations on value not supported
- **Distribution / Partitioning**

Now as these keys uniquely identify each data set,
it is much easier to distribute

doesn't have to be a schema, or in First Normal Form

Aside: Hash Functions

- A function that maps any data to a “hash value” (e.g., an integer)



Aside: Hash Functions

- Example: data and hash value are integers
- Simple hash function:
 - $h(\text{key}) = \text{key \% 42};$
 - $h(10) = 10$
 - $h(2) = 2$
 - $h(50) = 8$
- What does this have to do with data distribution?

Key-Value Stores Features

- **Data model:** (key,value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - `get(key)`, `put(key, value)`
 - Operations on value not supported USE HASH FUNCTIONS to decide what server to put on
- **Distribution / Partitioning** – w/ hash function
 - No replication: key k is stored at server $h(k)$
 - 3-way replication: key k stored at $h_1(k), h_2(k), h_3(k)$

How does `get(k)` work? How does `put(k,v)` work?

•²⁸

apply hash, and the output tells us what server data is on. Look there.

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?

How does query processing work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?
- Option 1: key=fid, value=entire flight record

How does query processing work?

•30

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?
- Option 1: key=fid, value=entire flight record
- Option 2: key=date, value=all flights that day

How does query processing work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)

Carriers(cid, name)

Example

- How would you represent the Flights data as key, value pairs?
have a bunch of choices to decide how you organize the data; carefully choosing key value pair can make certain operations much faster.
- Option 1: key=fid, value=entire flight record
- Option 2: key=date, value=all flights that day
- Option 3: key=(origin,dest), value=all flights between

How does query processing work?

Data Models

Taxonomy based on data models:

- **Key-value stores**
 - e.g., Project Voldemort, Memcached
- ☞ • **Document stores**
 - e.g., SimpleDB, CouchDB, MongoDB

Motivation

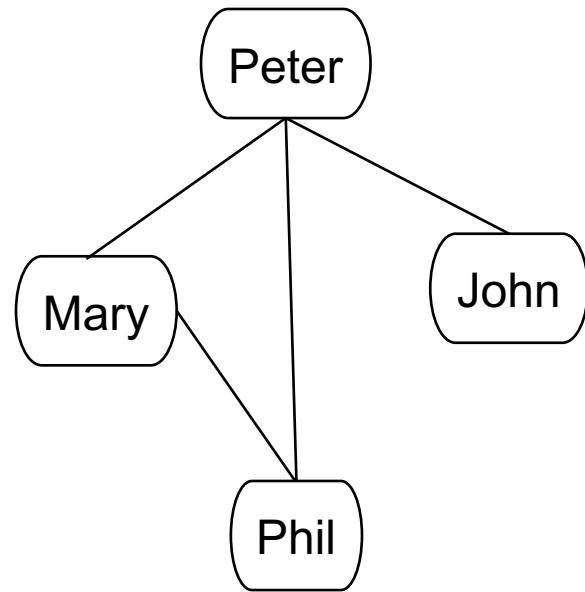
- In Key, Value stores, the Value is often a very complex object
 - Key = ‘2010/7/1’, Value = [all flights that date]
Document Store: Force Value to have a specific format
- Better: allow DBMS to understand the *value*
 - Represent *value* as a JSON (or XML...) document
 - [all flights on that date] = a JSON file
 - May search for all flights on a given date

Document Stores Features

- **Data model:** (key,document) pairs
 - Key = string/integer, unique for the entire data
 - Document = JSON, or XML
- **Operations**
 - Get/put document by key
 - Query language over JSON
- **Distribution / Partitioning**
 - Entire documents, as for key/value pairs

We will discuss JSON

Example: storing FB friends



As a graph

OR

Add new attributes?

Storing lists?

Person1	Person2	is_friend
Peter	John	1
John	Mary	0
Mary	Phil	1
Phil	Peter	1
...

As a relation

We will learn the tradeoffs of different data models later this quarter

JSON

JSON - Overview

- JavaScript Object Notation = lightweight text-based open standard designed for human-readable data interchange. Interfaces in C, C++, Java, Python, Perl, etc.
- The filename extension is .json.

We will emphasize JSON as semi-structured data

JSON Syntax

```
{ "book": [  
    {"id": "01",  
        "language": "Java",  
        "author": "H. Javeson",  
        "year": 2015  
    },  
    {"id": "07",  
        "language": "C++",  
        "edition": "second"  
        "author": "E. Sepp",  
        "price": 22.25  
    }  
]
```

JSon vs Relational

- Relational data model
 - Rigid flat structure (tables)
 - Schema must be fixed in advance
 - Binary representation: good for performance, bad for exchange
 - Query language based on Relational Calculus
- Semistructured data model / JSon
 - Flexible, nested structure (trees) do NOT have a fixed schema
 - Does not require predefined schema ("self describing")
 - Text representation: good for exchange, bad for performance human readable
 - Most common use: Language API; query languages emerging

JSon Terminology

- Data is represented in name/value pairs.
- Curly braces hold objects
 - Each object is a list of name/value pairs separated by ,(comma)
 - Each pair is a name followed by ':'(colon) followed by the value
- Square brackets hold arrays and values are separated by ,(comma).

JSon Data Structures

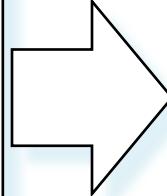
where value can be a complex object
(i.e. can store another table within a
value)

- Collections of name-value pairs:
 - {“name1”: value1, “name2”: value2, ...}
 - The “name” is also called a “key”
- Ordered lists of values:
 - [obj1, obj2, obj3, ...]

Avoid Using Duplicate Keys

The standard allows them, but many implementations don't

```
{"id": "07",      DONT DO THIS  
  "title": "Databases",  
  "author": "Garcia-Molina",  
  "author": "Ullman",  
  "author": "Widom"  
}
```



```
{"id": "07",  
  "title": "Databases",  
  "author": ["Garcia-Molina",  
             "Ullman",  
             "Widom"]}
```

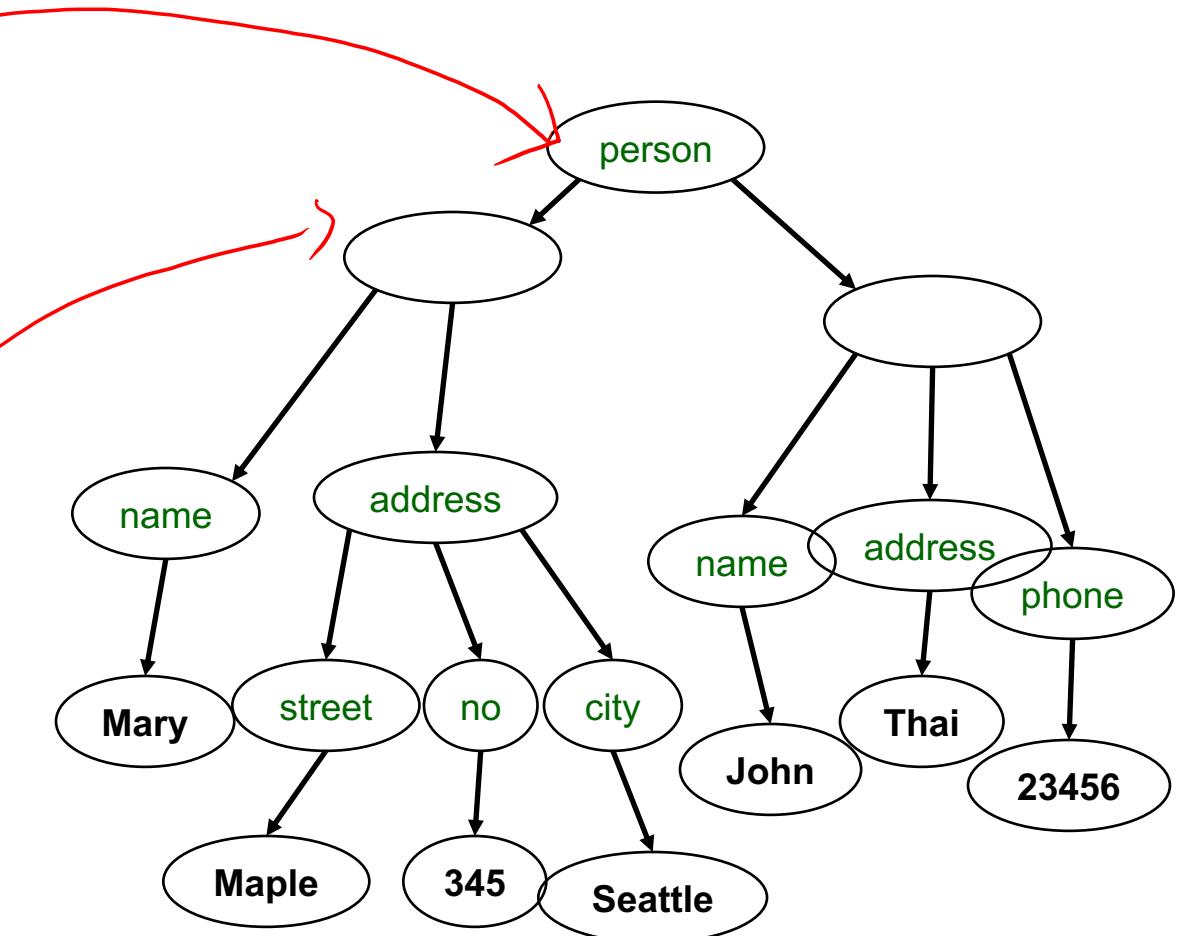
[1]

JSon Datatypes

- Number
- String = double-quoted
- Boolean = true or false
- nullempty

JSon Semantics: a Tree !

```
{"person":  
[ {"name": "Mary",  
"address":  
{"street": "Maple",  
"no": 345,  
"city": "Seattle"}},  
 {"name": "John",  
"address": "Thailand",  
"phone": 2345678}  
]  
}
```



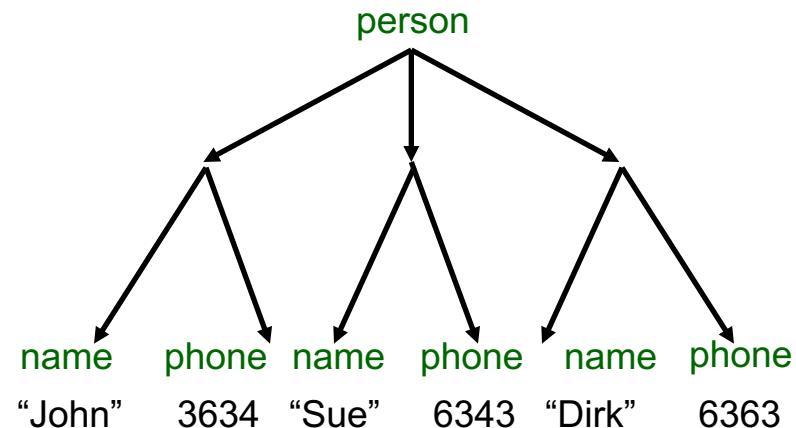
JSon Data

- JSon is **self-describing** NO SET SCHEMA; keys of JSON tells us what our attributes are
- Schema elements become part of the data
 - Relational schema: person(name,phone)
 - In Json “person”, “name”, “phone” are part of the data, and are repeated many times
- Consequence: JSon is much more flexible
- JSon = **semistructured** data
 - i.e. attributes can vary across tuples

Mapping Relational Data to JSON

Person

name	phone
John	3634
Sue	6343
Dirk	6363



```
{"person":  
  [{"name": "John", "phone":3634},  
   {"name": "Sue", "phone":6343},  
   {"name": "Dirk", "phone":6383}  
 ]  
}
```

Mapping Relational Data to JSON

May inline foreign keys

Person

name	phone
John	3634
Sue	6343

CAN STORE SUB-tables ->
don't need joins

Orders

personName	date	product
John	2002	Gizmo
John	2004	Gadget
Sue	2002	Gadget

```
{“Person”:  
[{"name": "John",  
“phone”:3646,  
“Orders”:[{"date":2002,  
“product”:"Gizmo"},  
 {"date":2004,  
“product”:"Gadget"}  
]  
},  
 {"name": "Sue",  
“phone”:6343,  
“Orders”:[{"date":2002,  
“product”:"Gadget"}  
]  
}]}
```