

Introduction to Data Management

CSE 414

Unit 6: Conceptual Design
E/R Diagrams
Integrity Constraints
BCNF

(3 lectures)

Introduction to Data Management

CSE 414

Integrity Constraints

Integrity Constraints Motivation

Primary Key is an example of an integrity constraint

fix fault data?

An integrity constraint is a condition specified on a database schema that restricts the data that can be stored in an instance of the database.

- ICs help prevent entry of incorrect information
- How? DBMS enforces integrity constraints
 - Allows only legal database instances (i.e., those that satisfy all constraints) to exist
 - Ensures that all necessary checks are always performed and avoids duplicating the verification logic in each application

balancing number of constraints (takes extra work to verify constraints are met) vs versatility to bad data

Constraints in E/R Diagrams

Finding constraints is part of the modeling process.
Commonly used constraints:

Keys: social security number uniquely identifies a person.

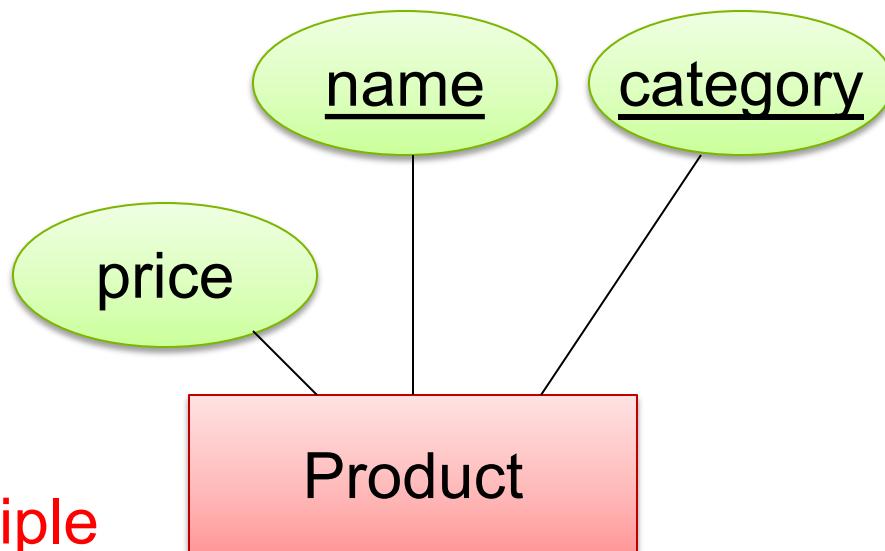
Single-value constraints: a person can have only one father.
i.e. many to one or one to one relationships are properly maintained!

Referential integrity constraints: if you work for a company, it must exist in the database.

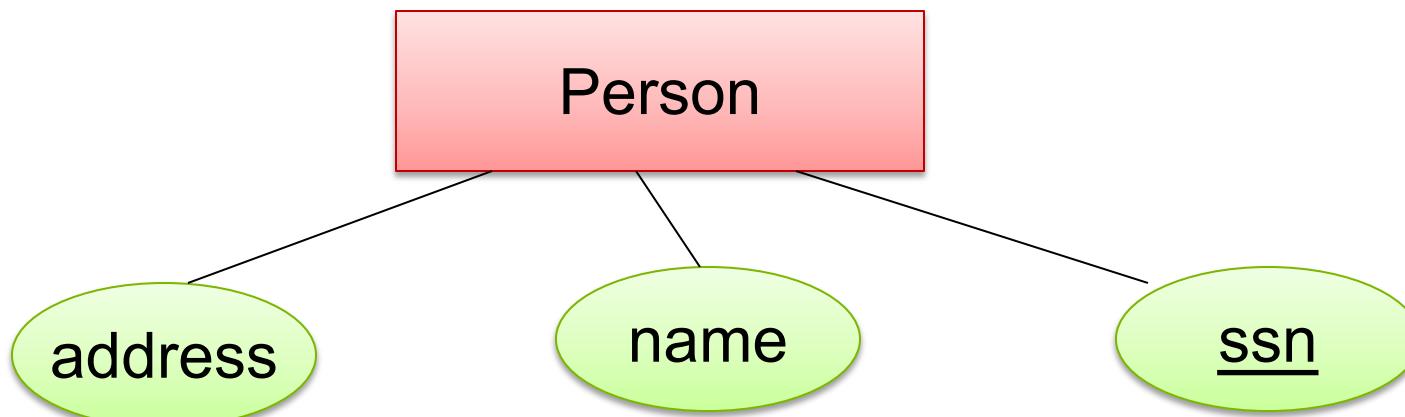
Other constraints: peoples' ages are between 0 and 150.

Keys in E/R Diagrams

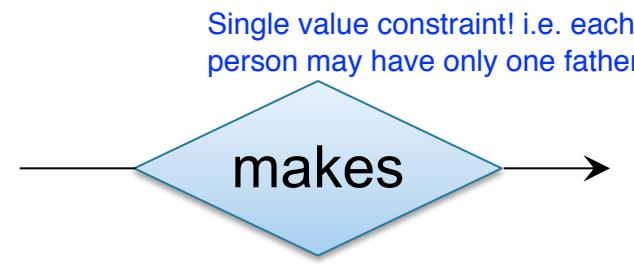
PRIMARY KEY
Underline:



No formal way
to specify multiple
keys in E/R diagrams



Single Value Constraints

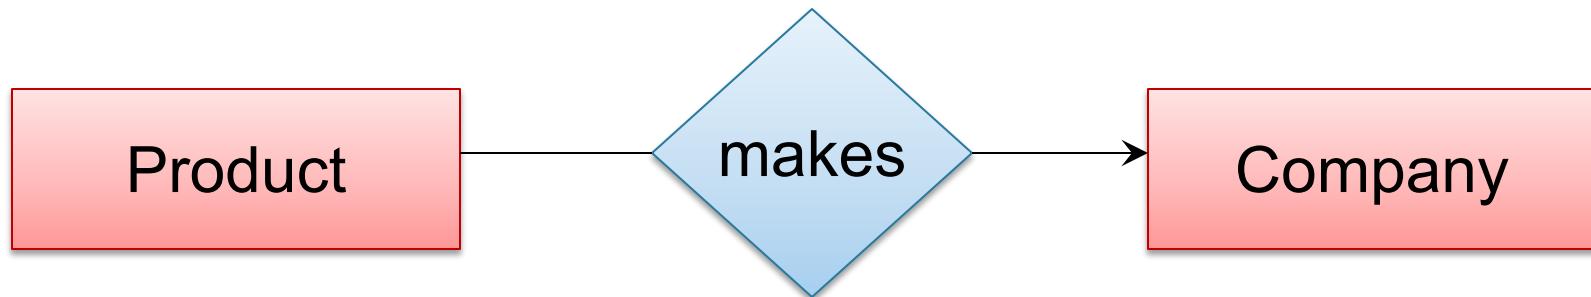


vs.



Referential Integrity Constraints

each product is made by at most one company,
but each company may make many products



Each product made by at most one company.
Some products made by no company

Each product made by exactly one company,
but each company may make many products

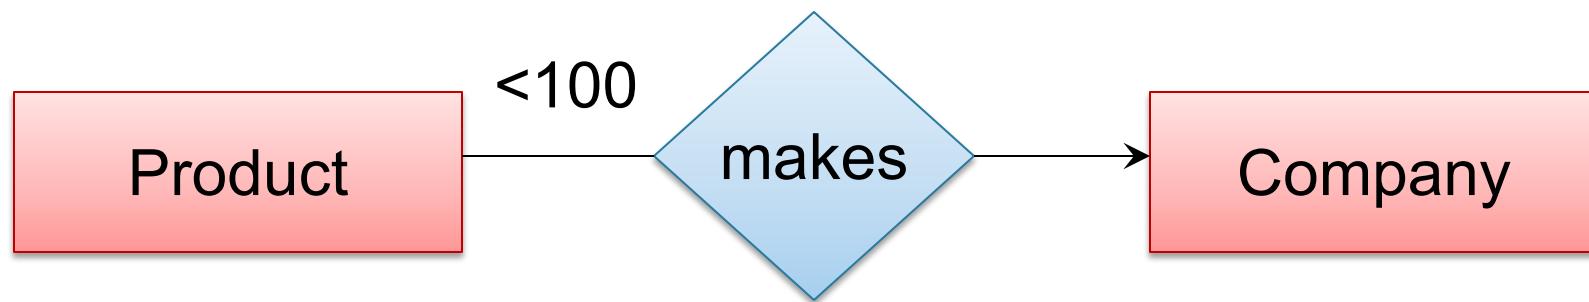


Each product made by exactly one company.

Other Constraints

Put other constraints on number of connections each element of one entity may have.

Many to one, but many < 100



Q: What does this mean ?

A: A Company entity cannot be connected by relationship to more than 99 Product entities

Constraints in SQL

Constraints in SQL:

- Keys, foreign keys
- Attribute-level constraints
- Tuple-level constraints
- Global constraints: assertions
 - multiple attributes?
- The more complex the constraint, the harder it is to check and to enforce

simplest

Most
complex

Key Constraints

Product(name, category)

```
CREATE TABLE Product (
    name CHAR(30) PRIMARY KEY,
    category VARCHAR(20))
```

OR:

```
CREATE TABLE Product (
    name CHAR(30),
    category VARCHAR(20),
    PRIMARY KEY (name))
```

puts a constraint on our data! cannot insert another tuple with same primary key!

Keys with Multiple Attributes

Product(name, category, price)

```
CREATE TABLE Product (
    name CHAR(30),
    category VARCHAR(20),
    price INT,
    PRIMARY KEY (name, category))
```

Name	Category	Price
Gizmo	Gadget	10
Camera	Photo	20
Gizmo	Photo	30
Gizmo	Gadget	40

Other Keys

```
CREATE TABLE Product (
    productID CHAR(10),
    name CHAR(30),
    category VARCHAR(20),
    price INT,
    PRIMARY KEY (productID),
    UNIQUE (name, category))
```

There is at most one **PRIMARY KEY**;
there can be many **UNIQUE**

i.e. other keys!
(name, category) is another key
name or category attribute values may repeat!
It is only the PAIR that forms a unique key

Foreign Key Constraints

foreign key must point to a key in another table (but this key does not necessarily need to be the Primary one)!

```
CREATE TABLE Purchase (
    prodName CHAR(30)
    REFERENCES Product(name),
    date DATETIME)
```

Referential
integrity
constraints

prodName is a **foreign key** to Product(name)
name must be a **key** in Product

May write
just Product
if name is PK

Foreign Key Constraints

- Example with multi-attribute primary key

```
CREATE TABLE Purchase (
    prodName CHAR(30),
    category VARCHAR(20),
    date DATETIME,
    FOREIGN KEY (prodName, category)
    REFERENCES Product(name, category)
```

- (name, category) must be a KEY in Product

What happens when data changes?

Types of updates:

- In Purchase: insert/update
- In Product: delete/update

i.e. ProdName is foreign key to Product.name, so if we insert a tuple into Purchase (Toy, Wiz) where the foreign key value does NOT match any of the primary key values in Product, we reach an inconsistent state

OR

Delete primary key in Product, so that a foreign key in Purchase no longer has a matching primary key in product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

How should we choose to maintain this internal integrity?
What should we do when you ask to perform an action that will violate
the internal constraints?

What happens when data changes?

- SQL has three policies for maintaining referential integrity:
- NO ACTION reject violating modifications (default)
- CASCADE after delete/update do delete/update make a change, delete corresponding values in other tables that are causing the issue
- SET NULL set foreign-key field to NULL when it now has no more matching corresponding key in another table
- SET DEFAULT set foreign-key field to default value
 - need to be declared with column, e.g.,
`CREATE TABLE Product (pid INT DEFAULT 42)`

more expensive with
more references/
foreign keys in
tables

Maintaining Referential Integrity

```
CREATE TABLE Purchase (
    prodName CHAR(30),
    category VARCHAR(20),
    date DATETIME,
    FOREIGN KEY (prodName, category)
        REFERENCES Product(name, category)
        ON UPDATE CASCADE
        ON DELETE SET NULL )
```

changing attribute values within tuples,
delete tuples that cause referential issues

The diagram illustrates a foreign key relationship between the **Product** and **Purchase** tables. A curved arrow points from the **Product** table to the **Purchase** table, indicating that the **ProdName** and **Category** columns in the **Purchase** table refer back to the **Name** and **Category** columns in the **Product** table.

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

ProdName	Category
Gizmo	Gizmo
Snap	Camera
EasyShoot	Camera

Constraints on Attributes and Tuples

- Constraints on attributes:

attribute cannot be null

check whether attribute
value satisfies specific
conditions

NOT NULL

CHECK condition

-- obvious meaning...
-- any condition !

- Constraints on tuples

CHECK condition

Constraints on Attributes and Tuples

```
CREATE TABLE R (
    A int NOT NULL,           attribute level constraint;
                                check if cond satisfied
    B int CHECK (B > 50 and B < 100),
    C varchar(20),
    D int,                   tuple level constraint! Most expensive!
                                CHECK (C >= 'd' or D > 0))
```

Constraints on Attributes and Tuples

```
CREATE TABLE Product (
    productID CHAR(10),
    name CHAR(30),
    category VARCHAR(20),
    price INT CHECK (price > 0),
    PRIMARY KEY (productID),
    UNIQUE (name, category))
```

Constraints on Attributes and Tuples

What does this constraint do?

```
CREATE TABLE Purchase (
    prodName CHAR(30)
        CHECK (prodName IN
            (SELECT Product.name
                FROM Product),
    date DATETIME NOT NULL)
```

Subqueries to check constraints!
in this case, the constraint is that if you insert
a tuple to the purchase table,
the associated product name for that purchase
must be in the product table!
VERY expensive. Every insert must run subquery

What
is the difference from
Foreign-Key ?

this does not require Product.name
to be unique/key in Product table!

General Assertions

checks entire SQL queries!

```
CREATE ASSERTION myAssert CHECK
(NOT EXISTS(
    SELECT Product.name
    FROM Product, Purchase
    WHERE Product.name = Purchase.prodName
    GROUP BY Product.name
    HAVING count(*) > 200) )
```

But most DBMSs do not implement assertions
Because it is hard to support them efficiently
Instead, they provide triggers

Introduction to Data Management

CSE 414

Design Theory and BCNF

Relational Schema Design

Name	<u>SSN</u>	<u>PhoneNumber</u>	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield

One person may have multiple phones, but lives in only one city

Primary key is thus (SSN, PhoneNumber)

What is the problem with this schema?

Relational Schema Design



Name	<u>SSN</u>	<u>PhoneNumber</u>	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield

Key should be SSN

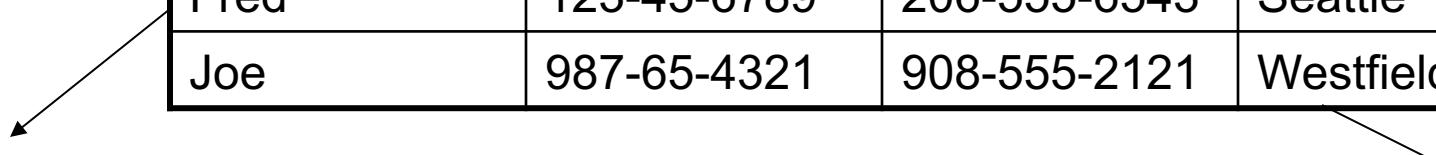
Anomalies:

- Redundancy = repeat data
- Update anomalies = what if Fred moves to “Bellevue”? might only update single tuple rather than both!
- Deletion anomalies = what if Joe deletes his phone number?

Relation Decomposition

Break the relation into two:

Name	SSN	PhoneNumber	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield



Name	<u>SSN</u>	City
Fred	123-45-6789	Seattle
Joe	987-65-4321	Westfield

<u>SSN</u>	PhoneNumber
123-45-6789	206-555-1234
123-45-6789	206-555-6543
987-65-4321	908-555-2121

Anomalies have gone:

- No more repeated data
- Easy to move Fred to “Bellevue” (how ?)
- Easy to delete all Joe’s phone numbers (how ?)

can delete full tuples!

Relational Schema Design (or Logical Design)

How do we do this systematically?

- Start with some relational schema
- Find out its *functional dependencies* (FDs)

- Use FDs to *normalize* the relational schema

Functional Dependencies (FDs)

Definition

If two tuples agree on the attributes

A_1, A_2, \dots, A_n

then they must also agree on the attributes

B_1, B_2, \dots, B_m

Formally:

$A_1 \dots A_n \xrightarrow{\text{determines}} B_1 \dots B_m$

$A_1, A_2, \dots, A_n \rightarrow B_1, B_2, \dots, B_m$

Functional Dependencies (FDs)

Definition $A_1, \dots, A_m \rightarrow B_1, \dots, B_n$ holds in R if:

$\forall t, t' \in R,$

$(t.A_1 = t'.A_1 \wedge \dots \wedge t.A_m = t'.A_m \rightarrow t.B_1 = t'.B_1 \wedge \dots \wedge t.B_n = t'.B_n)$

R	A_1	\dots	A_m		B_1	\dots	B_n	
t								
t'								

if t, t' agree here then t, t' agree here

Example

An FD holds, or does not hold on an instance:

EmplID	Name	Phone	Position
E0045	Smith	1234	Clerk
E3542	Mike	9876	Salesrep
E1111	Smith	9876	Salesrep
E9999	Mary	1234	Lawyer

EmplID → Name, Phone, Position

Position → Phone

but not Phone → Position

Example

EmplID	Name	Phone	Position
E0045	Smith	1234	Clerk
E3542	Mike	9876 ←	Salesrep
E1111	Smith	9876 ←	Salesrep
E9999	Mary	1234	Lawyer

Position → Phone

Example

EmplID	Name	Phone	Position
E0045	Smith	1234 →	Clerk
E3542	Mike	9876	Salesrep
E1111	Smith	9876	Salesrep
E9999	Mary	1234 →	Lawyer

But not Phone → Position

Example

name → color
category → department
color, category → price

name	category	color	department	price
Gizmo	Gadget	Green	Toys	49
Tweaker	Gadget	Green	Toys	99

Do all the FDs hold on this instance?

Example

name → color
category → department
color, category → price

name	category	color	department	price
Gizmo	Gadget	Green	Toys	49
Tweaker	Gadget	Green	Toys	49
Gizmo	Stationary	Green	Office-supp.	59

What about this one ?

CSE 414 – Autumn 2018

76

Buzzwords

- FD **holds** or **does not hold** on an instance
- If we can be sure that *every instance* of R will be one in which a given FD is true, then we say that **R satisfies the FD**
- If we say that R satisfies an FD, we are **stating a constraint on R**

An Interesting Observation

If all these FDs are true:

$\text{name} \rightarrow \text{color}$
 $\text{category} \rightarrow \text{department}$
 $\text{color, category} \rightarrow \text{price}$

Then this FD also holds:

$\text{name, category} \rightarrow \text{price}$

An Interesting Observation

If all these FDs are true:

$\text{name} \rightarrow \text{color}$
 $\text{category} \rightarrow \text{department}$
 $\text{color, category} \rightarrow \text{price}$

Then this FD also holds:

$\text{name, category} \rightarrow \text{price}$

An Interesting Observation

If all these FDs are true:

$\text{name} \rightarrow \text{color}$
 $\text{category} \rightarrow \text{department}$
 $\text{color, category} \rightarrow \text{price}$

Then this FD also holds:

$\text{name, category} \rightarrow \text{price}$

If we find out from application domain that a relation satisfies some FDs, it doesn't mean that we found all the FDs that it satisfies! There could be more FDs implied by the ones we have.

Closure of a set of Attributes

Given a set of attributes A_1, \dots, A_n

The **closure** is the set of attributes B, notated $\{A_1, \dots, A_n\}^+$,
s.t. $A_1, \dots, A_n \rightarrow B$

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

Closures:

$$\text{name}^+ = \{\text{name, color}\}$$

$$\text{color}^+ = \{\text{color}\}$$

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$$\{\text{name, category}\}^+ = \{ \text{ name, category, } \}$$

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$
 $\{ \text{ name, category, color, } \}$

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$
 $\{ \text{ name, category, color, department } \}$

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$
 { name, category, color, department, price }

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$
 $\{ \text{name, category, color, department, price} \}$

Hence: $\text{name, category} \rightarrow \text{color, department, price}$

Example

In class:

$R(A, B, C, D, E, F)$

$A, B \rightarrow C$
$A, D \rightarrow E$
$B \rightarrow D$
$A, F \rightarrow B$

Compute $\{A, B\}^+$ $X = \{A, B,$ }

Compute $\{A, F\}^+$ $X = \{A, F,$ }

Example

In class:

$R(A, B, C, D, E, F)$

$A, B \rightarrow C$
$A, D \rightarrow E$
$B \rightarrow D$
$A, F \rightarrow B$

Compute $\{A, B\}^+$ $X = \{A, B, C, D, E\}$

Compute $\{A, F\}^+$ $X = \{A, F, \}$

Example

In class:

$R(A, B, C, D, E, F)$

$A, B \rightarrow C$
$A, D \rightarrow E$
$B \rightarrow D$
$A, F \rightarrow B$

Compute $\{A, B\}^+$ $X = \{A, B, C, D, E\}$

Compute $\{A, F\}^+$ $X = \{A, F, B, C, D, E\}$

Example

In class:

$R(A,B,C,D,E,F)$

A, B → C
A, D → E
B → D
A, F → B

Compute $\{A, B\}^+$ $X = \{A, B, C, D, E\}$

Compute $\{A, F\}^+$ $X = \{A, F, B, C, D, E\}$

Practice at Home

Find all FD's implied by:

$$\begin{array}{l} A, B \rightarrow C \\ A, D \rightarrow B \\ B \rightarrow D \end{array}$$

Practice at Home

Find all FD's implied by:

$$\begin{array}{l} A, B \rightarrow C \\ A, D \rightarrow B \\ B \rightarrow D \end{array}$$

Step 1: Compute X^+ , for every X :

$$A^+ = A, \quad B^+ = BD, \quad C^+ = C, \quad D^+ = D$$

$$\begin{aligned} AB^+ &= ABCD, \quad AC^+ = AC, \quad AD^+ = ABCD, \\ &\quad BC^+ = BCD, \quad BD^+ = BD, \quad CD^+ = CD \end{aligned}$$

$$ABC^+ = ABD^+ = ACD^+ = ABCD \text{ (no need to compute— why ?)}$$

$$BCD^+ = BCD, \quad ABCD^+ = ABCD$$

Step 2: Enumerate all FD's $X \rightarrow Y$, s.t. $Y \subseteq X^+$ and $X \cap Y = \emptyset$:

$$AB \rightarrow CD, \quad AD \rightarrow BC, \quad ABC \rightarrow D, \quad ABD \rightarrow C, \quad ACD \rightarrow B$$