

Introduction to Data Management

CSE 414

Unit 6: Conceptual Design
E/R Diagrams
Integrity Constraints
BCNF

(3 lectures)

Introduction to Data Management

CSE 414

Design Theory and BCNF

Relational Schema Design

Name	<u>SSN</u>	<u>PhoneNumber</u>	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield

One person may have multiple phones, but lives in only one city

Primary key is thus (SSN, PhoneNumber)

What is the problem with this schema?

Relational Schema Design

Name	SSN	PhoneNumber	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield

Anomalies:

- Redundancy = repeat data takes up extra space in our database unnecessarily!
- Update anomalies = what if Fred moves to “Bellevue”?
- Deletion anomalies = what if Joe deletes his phone number?
how do we handle this? delete whole tuple? leave null value?

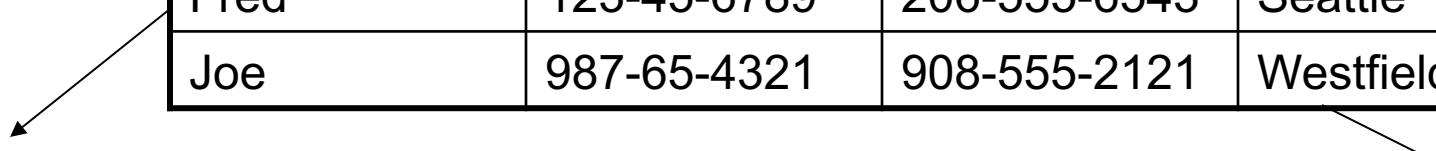
repeated data makes updates more prone to error; need to change both entries!

Relation Decomposition

Break the relation into two:

SSN should be unique for each person!

Name	SSN	PhoneNumber	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield



Name	<u>SSN</u>	City
Fred	123-45-6789	Seattle
Joe	987-65-4321	Westfield

<u>SSN</u>	PhoneNumber
123-45-6789	206-555-1234
123-45-6789	206-555-6543
987-65-4321	908-555-2121

Anomalies have gone:

- No more repeated data
- Easy to move Fred to “Bellevue” (how ?)
- Easy to delete all Joe’s phone numbers (how ?)

change single City entry associated to
FRED

Relational Schema Design (or Logical Design)

How do we do this systematically?

- Start with some relational schema
- Find out its *functional dependencies* (FDs)
GENERALIZED concept of a Key
- Use FDs to *normalize* the relational schema

Functional Dependencies (FDs)

Definition

If two tuples agree on the attributes

A_1, A_2, \dots, A_n

then they must also agree on the attributes

B_1, B_2, \dots, B_m

Formally:

$A_1 \dots A_n$ determines $B_1 \dots B_m$

$A_1, A_2, \dots, A_n \rightarrow B_1, B_2, \dots, B_m$

Functional Dependencies (FDs)

Definition $A_1, \dots, A_m \rightarrow B_1, \dots, B_n$ holds in R if:

$\forall t, t' \in R,$

$(t.A_1 = t'.A_1 \wedge \dots \wedge t.A_m = t'.A_m \rightarrow t.B_1 = t'.B_1 \wedge \dots \wedge t.B_n = t'.B_n)$

R	A_1	\dots	A_m		B_1	\dots	B_n	
t								
t'								

if t, t' agree here then t, t' agree here

subset of attributes!

another subset of attributes!

Example

An FD holds, or does not hold on an instance:

EmplID	Name	Phone	Position
E0045	Smith	1234	Clerk
E3542	Mike	9876	Salesrep
E1111	Smith	9876	Salesrep
E9999	Mary	1234	Lawyer

EmplID → Name, Phone, Position

Position → Phone

but not Phone → Position

Example

EmplID	Name	Phone	Position
E0045	Smith	1234	Clerk
E3542	Mike	9876 ←	Salesrep
E1111	Smith	9876 ←	Salesrep
E9999	Mary	1234	Lawyer

Position → Phone

Example

EmplID	Name	Phone	Position
E0045	Smith	1234 →	Clerk
E3542	Mike	9876	Salesrep
E1111	Smith	9876	Salesrep
E9999	Mary	1234 →	Lawyer

But not Phone → Position

Example

holds
DOES NOT HOLD

name \rightarrow color
category \rightarrow department
color, category \rightarrow price

name	category	color	department	price
Gizmo	Gadget	Green	Toys	49
Tweaker	Gadget	Green	Toys	99

Do all the FDs hold on this instance?

Example

holds

holds

holds

name → color
category → department
color, category → price

name	category	color	department	price
Gizmo	Gadget	Green	Toys	49
Tweaker	Gadget	Green	Toys	49
Gizmo	Stationary	Green	Office-supp.	59

What about this one ?

CSE 414 – Autumn 2018

83

Buzzwords

- FD **holds** or **does not hold** on an instance
 - Places a constraint on our table
- If we can be sure that every instance of R will be one in which a given FD is true, then we say that **R satisfies the FD**
- If we say that R satisfies an FD, we are **stating a constraint on R**

Why bother with FDs?

Name	<u>SSN</u>	<u>PhoneNumber</u>	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield

Anomalies:

- Redundancy = repeat data
- Update anomalies = what if Fred moves to “Bellevue”?
- Deletion anomalies = what if Joe deletes his phone number?

An Interesting Observation

If all these FDs are true:

$\text{name} \rightarrow \text{color}$
 $\text{category} \rightarrow \text{department}$
 $\text{color, category} \rightarrow \text{price}$

can derive further Functional Dependencies using relationships

$A \rightarrow B \rightarrow C$ implies $A \rightarrow C$

Then this FD also holds:

$\text{name, category} \rightarrow \text{price}$

An Interesting Observation

If all these FDs are true:

$\text{name} \rightarrow \text{color}$
 $\text{category} \rightarrow \text{department}$
 $\text{color, category} \rightarrow \text{price}$

Then this FD also holds:

$\text{name, category} \rightarrow \text{price}$

If we find out from application domain that a relation satisfies some FDs, it doesn't mean that we found all the FDs that it satisfies! There could be more FDs implied by the ones we have.

Closure of a set of Attributes

Given a set of attributes A_1, \dots, A_n

The **closure** is the set of attributes B, notated $\{A_1, \dots, A_n\}^+$,
s.t. $A_1, \dots, A_n \rightarrow B$

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

Closures:

$$\text{name}^+ = \{\text{name, color}\}$$

everything that you can derive from functional dependencies!

$$\text{color}^+ = \{\text{color}\}$$

Note that each attribute trivially determines itself!

NO CLOSURE SET can be empty because of this!

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$$\{\text{name, category}\}^+ = \{ \text{ name, category, } \}$$

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$
 $\{ \text{ name, category}, \text{color}, \text{ }$ name determines color $\}$

Closure Algorithm

$$X = \{A_1, \dots, A_n\}.$$

Repeat until X doesn't change do:

if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X.

Example:

1. name → color
 2. category → department
 3. color, category → price

$\{\text{name, category}\}^+ = \{ \text{name, category, color, department} \}$

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$

{ name, category, color, department, price }

color and category determine price

de

Closure Algorithm

$X = \{A_1, \dots, A_n\}$.

Repeat until X doesn't change **do**:
if $B_1, \dots, B_n \rightarrow C$ is a FD **and**
 B_1, \dots, B_n are all in X
then add C to X .

Example:

1. name \rightarrow color
2. category \rightarrow department
3. color, category \rightarrow price

$\{\text{name, category}\}^+ =$
 $\{ \text{name, category, color, department, price} \}$

Hence: $\text{name, category} \rightarrow \text{color, department, price}$

Example

In class:

$R(A, B, C, D, E, F)$

$A, B \rightarrow C$
$A, D \rightarrow E$
$B \rightarrow D$
$A, F \rightarrow B$

Compute $\{A, B\}^+$ $X = \{A, B,$ }

Compute $\{A, F\}^+$ $X = \{A, F,$ }

Example

In class:

$R(A, B, C, D, E, F)$

$A, B \rightarrow C$
$A, D \rightarrow E$
$B \rightarrow D$
$A, F \rightarrow B$

Compute $\{A, B\}^+$ $X = \{A, B, C, D, E\}$ }

Compute $\{A, F\}^+$ $X = \{A, F\}$ }

Example

In class:

$R(A, B, C, D, E, F)$

$A, B \rightarrow C$
$A, D \rightarrow E$
$B \rightarrow D$
$A, F \rightarrow B$

Compute $\{A, B\}^+$ $X = \{A, B, C, D, E\}$ }

Compute $\{A, F\}^+$ $X = \{A, F, B, C, D, E\}$ }

Example

In class:

$R(A,B,C,D,E,F)$

A, B → C
A, D → E
B → D
A, F → B

Compute $\{A, B\}^+$ $X = \{A, B, C, D, E\}$

Compute $\{A, F\}^+$ $X = \{A, F, B, C, D, E\}$

Practice at Home

Find all FD's implied by:

$$\begin{array}{l} A, B \rightarrow C \\ A, D \rightarrow B \\ B \rightarrow D \end{array}$$

Practice at Home

Find all FD's implied by:

$$\begin{array}{l} A, B \rightarrow C \\ A, D \rightarrow B \\ B \rightarrow D \end{array}$$

Step 1: Compute X^+ , for every X :

$$A^+ = A, \quad B^+ = BD, \quad C^+ = C, \quad D^+ = D$$

$$\begin{aligned} AB^+ &= ABCD, \quad AC^+ = AC, \quad AD^+ = ABCD, \\ &\quad BC^+ = BCD, \quad BD^+ = BD, \quad CD^+ = CD \end{aligned}$$

$$ABC^+ = ABD^+ = ACD^+ = ABCD \text{ (no need to compute— why ?)}$$

$$BCD^+ = BCD, \quad ABCD^+ = ABCD$$

Step 2: Enumerate all FD's $X \rightarrow Y$, s.t. $Y \subseteq X^+$ and $X \cap Y = \emptyset$:

$$AB \rightarrow CD, \quad AD \rightarrow BC, \quad ABC \rightarrow D, \quad ABD \rightarrow C, \quad ACD \rightarrow B$$

100

Keys

- A **superkey** is a set of attributes A_1, \dots, A_n s.t. for any other attribute B , we have $A_1, \dots, A_n \rightarrow B$
- A **key** is a minimal superkey
 - A superkey and for which no subset is a superkey

smallest number of attributes (smallest subset of attributes)
which is a super key

Computing (Super)Keys

- For all sets X , compute X^+
- If $X^+ = [\text{all attributes}]$, then X is a superkey
- Try reducing to the minimal X 's to get the key

Example

Product(name, price, category, color)

name, category → price
category → color

What is the key ?

Example

Product(name, price, category, color)

name, category → price
category → color

What is the key ?

closure that contains all attributes



(name, category) + = { name, category, price, color }

of the smallest size = minimal superkey = key

Hence (name, category) is a key

Example

Product(name, price, category, color)

name, category → price
category → color

What is the key ?

(name, category) + = { name, category, price, color }

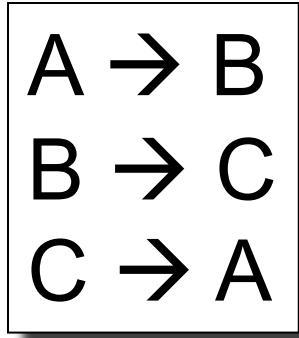
Key or Keys ?

Can we have more than one key ?

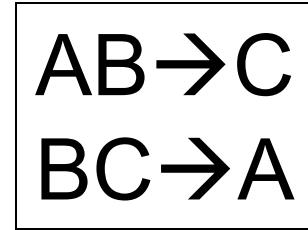
YES!

Given $R(A,B,C)$ define FD's s.t. there are two or more distinct keys

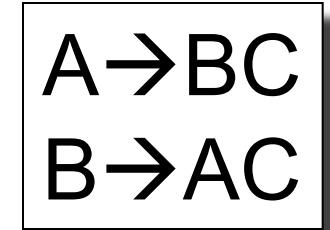
AB is a key as $\{AB\}^+ = \{A, B, C\}$
BC is a key



or



or



KEYS
A as $\{A\}^+ = \{A, B, C\}$
B
C

what are the keys here ?

Eliminating Anomalies

Main idea:

- $X \rightarrow A$ is OK if X is a (super)key
- $X \rightarrow A$ is not OK otherwise
 - Need to decompose the table, but how?

Boyce-Codd Normal Form

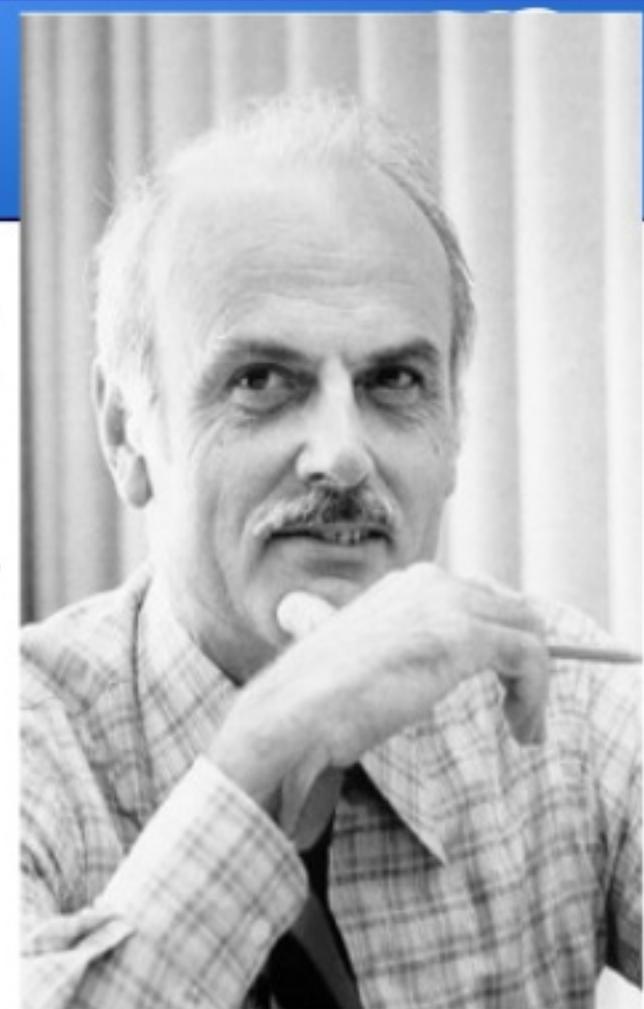
normalization is putting data table into Boyce-Codd Normal Form!

Boyce-Codd Normal Form

Dr. Raymond F. Boyce

Edgar Frank “Ted” Codd

"A Relational Model of Data for Large Shared Data Banks"



Boyce-Codd Normal Form

If there are no
“bad” FDs:

every functional dependency determines
all elements in table

Definition. A relation R is in BCNF if:

Whenever $X \rightarrow B$ is a non-trivial dependency,
then X is a superkey.

X determines everything directly or indirectly

Definition. A relation R is in BCNF if:

Equivalently: $\forall X$, either $X^+ = X$ (i.e., X is not in any FDs)

X is a trivial functional dependency ; $X \rightarrow X$
or $X^+ = [\text{all attributes}]$ (computed using FDs)

convert something not in BCNF into BCNF

BCNF Decomposition Algorithm

Normalize(R)

find a nontrivial closure set

find X such that $X \rightarrow B$ is non trivial functional dependency (B is not in X)

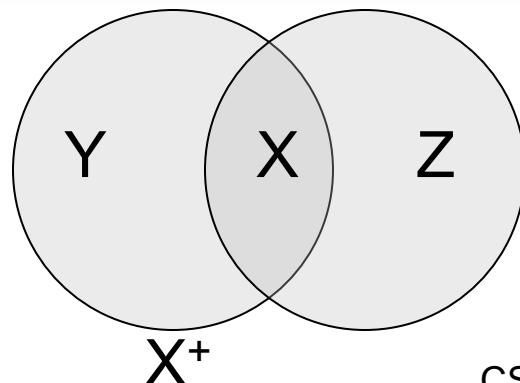
find X s.t.: $X \neq X^+$ and $X^+ \neq [all\ attributes]$

if (not found) **then** “R is in BCNF”

let $Y = X^+ - X$; $Z = [all\ attributes] - X^+$

decompose R into $R_1(X \cup Y)$ and $R_2(X \cup Z)$

Normalize(R_1); Normalize(R_2);

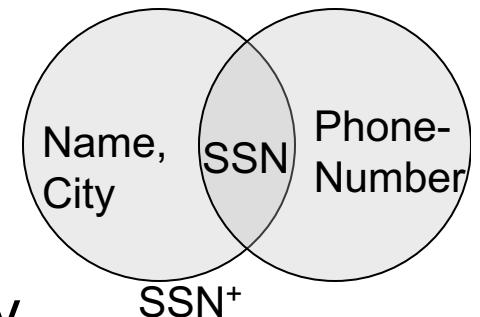


Example

Name	SSN	PhoneNumber	City
Fred	123-45-6789	206-555-1234	Seattle
Fred	123-45-6789	206-555-6543	Seattle
Joe	987-65-4321	908-555-2121	Westfield
Joe	987-65-4321	908-555-1234	Westfield

$\text{SSN} \rightarrow \text{Name, City}$

SSN IS NOT a super key as it does not determine phone number!



The only key is: {SSN, PhoneNumber}

Hence $\text{SSN} \rightarrow \text{Name, City}$ is a “bad” dependency

In other words:

$\text{SSN}^+ = \text{SSN, Name, City}$ and is neither SSN nor All Attributes¹¹³

not trivial but also not a super key

Example BCNF Decomposition

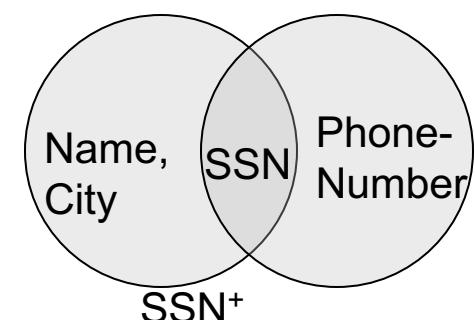
Name	<u>SSN</u>	City
Fred	123-45-6789	Seattle
Joe	987-65-4321	Westfield

This initial functional dependency became its own table (so that the table is BCNF)

$\text{SSN} \rightarrow \text{Name, City}$

<u>SSN</u>	<u>PhoneNumber</u>
123-45-6789	206-555-1234
123-45-6789	206-555-6543
987-65-4321	908-555-2121
987-65-4321	908-555-1234

unavoidable redundancy!



Let's check anomalies:

- Redundancy ?
- Update ?
- Delete ?

Find X s.t.: $X \neq X^+$ and $X^+ \neq [all\ attributes]$

Example BCNF Decomposition

Person(name, SSN, age, hairColor, phoneNumber)

$SSN \rightarrow name, age$

$age \rightarrow hairColor$

Find X s.t.: $X \neq X^+$ and $X^+ \neq [\text{all attributes}]$

Example BCNF Decomposition

$\text{Person}(\text{name}, \text{SSN}, \text{age}, \text{hairColor}, \text{phoneNumber})$

$\text{SSN} \rightarrow \text{name}, \text{age}$

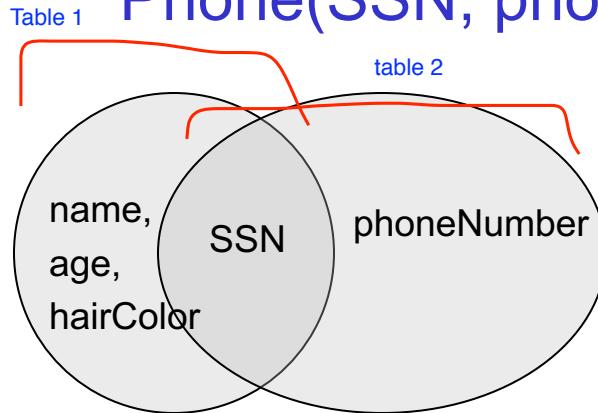
$\text{age} \rightarrow \text{hairColor}$

SSN does NOT determine every attribute!
"bad" functional dependency so we must split into multiple tables

Iteration 1: Person : $\text{SSN}^+ = \text{SSN}, \text{name}, \text{age}, \text{hairColor}$

Decompose into: $P(\text{SSN}, \text{name}, \text{age}, \text{hairColor})$

$\text{Phone}(\text{SSN}, \text{phoneNumber})$



Find X s.t.: $X \neq X^+$ and $X^+ \neq [\text{all attributes}]$

Example BCNF Decomposition

Person(name, SSN, age, hairColor, phoneNumber)

$\text{SSN} \rightarrow \text{name, age}$

$\text{age} \rightarrow \text{hairColor}$

What are
the keys ?

Iteration 1: Person: $\text{SSN}^+ = \text{SSN, name, age, hairColor}$

Decompose into: P(SSN, name, age, hairColor)

Phone(SSN, phoneNumber)

Iteration 2: P: $\text{age}^+ = \text{age, hairColor}$

Decompose: People(SSN, name, age)

Hair(age, hairColor)

Phone(SSN, phoneNumber)

age does NOT determine every
attribute in first table
bad functional dependency so
continue iterating!

now SSN and age determine all other fields;
super keys
run algorithm until there are no more divisions
that can occur

Find X s.t.: $X \neq X^+$ and $X^+ \neq [\text{all attributes}]$

Example BCNF Decomposition

$\text{Person}(\text{name}, \text{SSN}, \text{age}, \text{hairColor}, \text{phoneNumber})$

$\text{SSN} \rightarrow \text{name}, \text{age}$

$\text{age} \rightarrow \text{hairColor}$

Note the keys!

Iteration 1: Person : $\text{SSN}^+ = \text{SSN}, \text{name}, \text{age}, \text{hairColor}$

Decompose into: $P(\underline{\text{SSN}}, \text{name}, \text{age}, \text{hairColor})$

$\text{Phone}(\underline{\text{SSN}}, \text{phoneNumber})$

Iteration 2: P : $\text{age}^+ = \text{age}, \text{hairColor}$

Decompose: $\text{People}(\underline{\text{SSN}}, \text{name}, \text{age})$

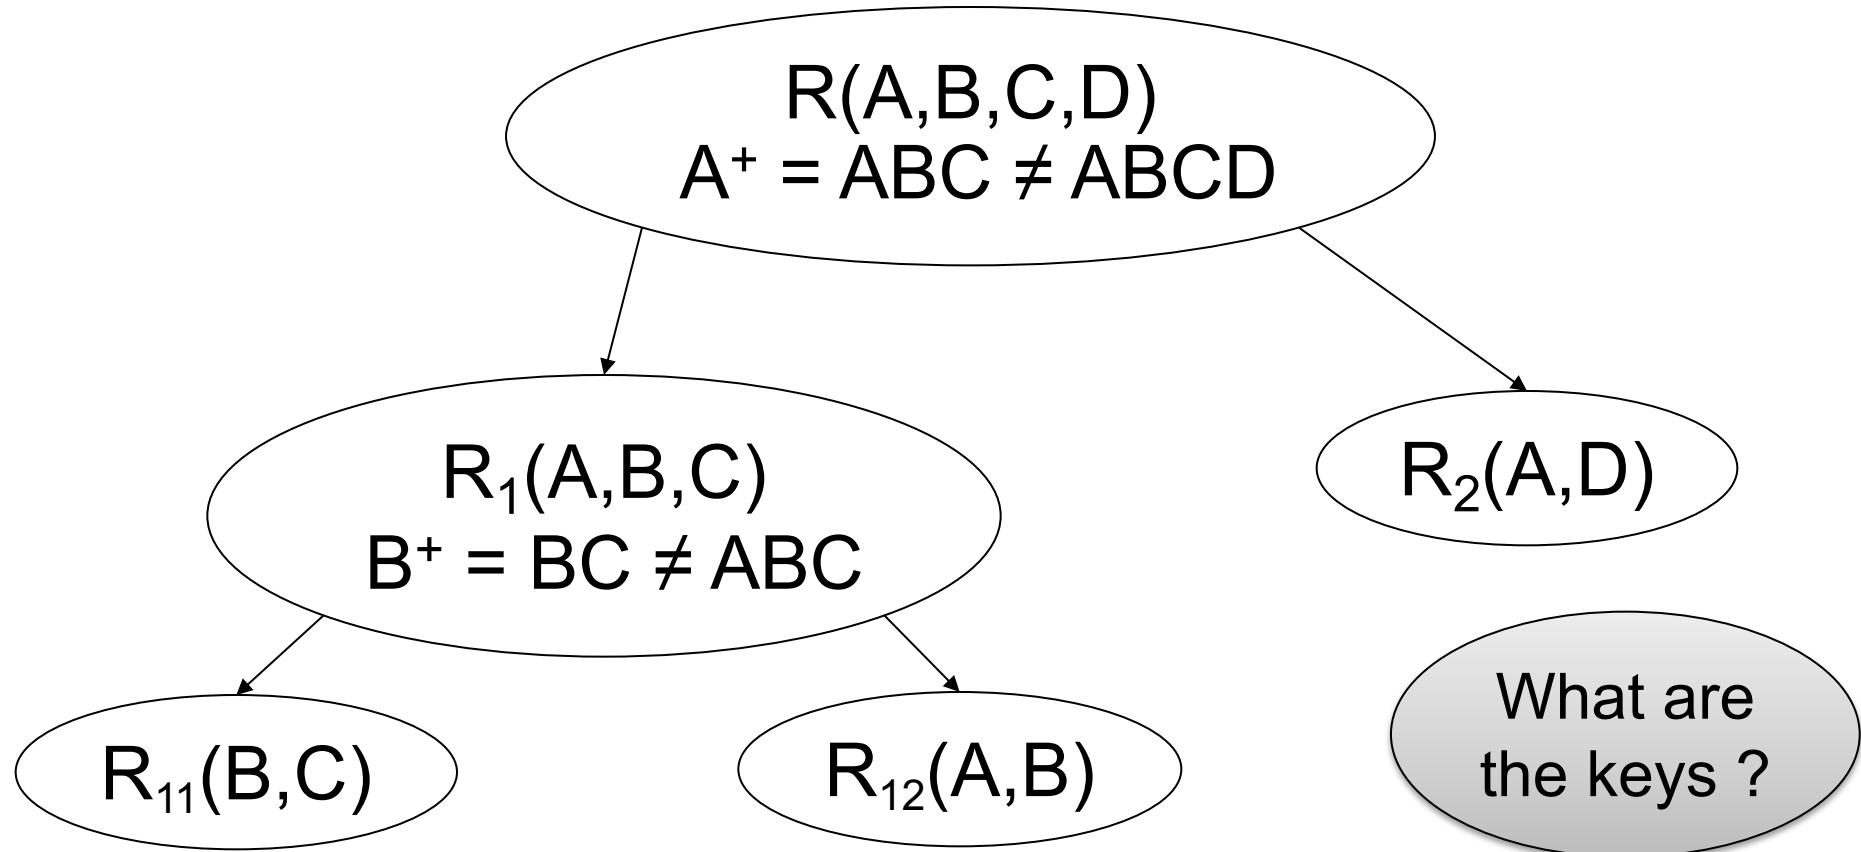
$\text{Hair}(\underline{\text{age}}, \text{hairColor})$

$\text{Phone}(\underline{\text{SSN}}, \text{phoneNumber})$

$R(A,B,C,D)$

$A \rightarrow B$
$B \rightarrow C$

Example: BCNF



What happens if in R we first pick B^+ ? Or AB^+ ?

B and A are the keys

Getting Practical

How to implement normalization in SQL

Motivation

- We learned about how to normalize tables to avoid anomalies
- How can we implement normalization in SQL if we can't modify existing tables?
 - This might be due to legacy applications that rely on previous schemas to run

Views

- A **view** in SQL =
 - A table computed from other tables, s.t., whenever the base tables are updated, the view is updated too
- More generally:
 - A **view** is derived data that keeps track of changes in the original data
- Compare:
 - A **function** computes a value from other values, but does not keep track of changes to the inputs

Purchase(customer, product, store)
Product(pname, price)

StorePrice(store, price)

A Simple View

Create a view that returns for each store
the prices of products purchased at that store

```
CREATE VIEW StorePrice AS
    SELECT DISTINCT x.store, y.price
    FROM Purchase x, Product y
    WHERE x.product = y.pname
```

This is like a new table
StorePrice(store,price)

Purchase(customer, product, store)
Product(pname, price)

StorePrice(store, price)

We Use a View Like Any Table

- A "high end" store is a store that sell some products over 1000.
- For each customer, return all the high end stores that they visit.

```
SELECT DISTINCT u.customer, u.store  
FROM Purchase u, StorePrice v  
WHERE u.store = v.store  
AND v.price > 1000
```

Types of Views

How are these different than a subquery?

- Virtual views pretty much a subquery
 - Computed only on-demand – slow at runtime
 - Always up to date
- Materialized views as soon as you define VIEW (CREATE VIEW)
it runs the query and essentially creates a new table
 - Pre-computed offline – fast at runtime
 - May have stale data (must recompute or update)
 - Indexes are materialized views
NOT up to date; must update periodically
- A key component of physical tuning of databases is the selection of materialized views and indexes

How view make things faster!

Vertical Partitioning

split by columns

Resumes

<u>SSN</u>	Name	Address	Resume	Picture
234234	Mary	Houston	Doc1...	JPG1...
345345	Sue	Seattle	Doc2...	JPG2...
345343	Joan	Seattle	Doc3...	JPG3...
432432	Ann	Portland	Doc4...	JPG4...

T1

<u>SSN</u>	Name	Address
234234	Mary	Houston
345345	Sue	Seattle
...		

T2

<u>SSN</u>	Resume
234234	Doc1...
345345	Doc2...

T3

<u>SSN</u>	Picture
234234	JPG1...
345345	JPG2...

T2.SSN is a key and a foreign key to T1.SSN. Same for T3.SSN 146

T1(ssn,name,address)
T2(ssn,resume)
T3(ssn,picture)

Resumes(ssn,name,address,resume,picture)

Vertical Partitioning

```
CREATE VIEW Resumes AS
    SELECT T1.ssn, T1.name, T1.address,
           T2.resume, T3.picture
    FROM   T1,T2,T3
    WHERE  T1.ssn=T2.ssn AND T1.ssn=T3.ssn
```

T1(ssn,name,address)
T2(ssn,resume)
T3(ssn,picture)

Resumes(ssn,name,address,resume,picture)

Vertical Partitioning

```
CREATE VIEW Resumes AS
    SELECT T1.ssn, T1.name, T1.address,
           T2.resume, T3.picture
    FROM   T1,T2,T3
    WHERE  T1.ssn=T2.ssn AND T1.ssn=T3.ssn
```

```
SELECT address
FROM   Resumes
WHERE  name = 'Sue'
```

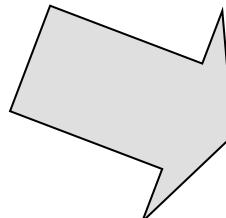
T1(ssn,name,address)
T2(ssn,resume)
T3(ssn,picture)

Resumes(ssn,name,address,resume,picture)

Vertical Partitioning

```
CREATE VIEW Resumes AS  
    SELECT T1.ssn, T1.name, T1.address,  
          T2.resume, T3.picture  
    FROM T1,T2,T3  
   WHERE T1.ssn=T2.ssn AND T1.ssn=T3.ssn
```

```
SELECT address  
FROM Resumes  
WHERE name = 'Sue'
```



Original query:

```
SELECT T1.address  
FROM T1, T2, T3  
WHERE T1.name = 'Sue'  
      AND T1.SSN=T2.SSN  
      AND T1.SSN = T3.SSN
```

T1(ssn,name,address)
T2(ssn,resume)
T3(ssn,picture)

Resumes(ssn,name,address,resume,picture)

Vertical Partitioning

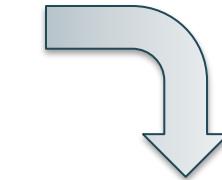
```
CREATE VIEW Resumes AS  
    SELECT T1.ssn, T1.name, T1.address,  
           T2.resume, T3.picture  
    FROM   T1,T2,T3  
    WHERE  T1.ssn=T2.ssn AND T1.ssn=T3.ssn
```

```
SELECT address  
FROM   Resumes  
WHERE  name = 'Sue'
```

Final query:

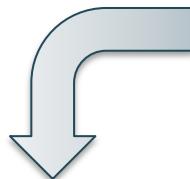
```
SELECT T1.address  
FROM T1  
WHERE T1.name = 'Sue'
```

database can automatically
filter out the unnecessary
calls (joins)



Modified query:

```
SELECT T1.address  
FROM T1, T2, T3  
WHERE T1.name = 'Sue'  
      AND T1.SSN=T2.SSN  
      AND T1.SSN = T3.SSN
```



Vertical Partitioning Applications

- **Advantages**
 - Speeds up queries that touch only a small fraction of columns
 - Single column can be compressed effectively, reducing disk I/O
- **Disadvantages**
 - Updates are expensive!
 - Need many joins to access many columns
 - Repeated key columns add overhead

have to edit many tables
at the same time

Horizontal Partitioning

Customers

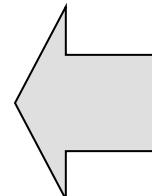
SSN	Name	City
234234	Mary	Houston
345345	Sue	Seattle
345343	Joan	Seattle
234234	Ann	Portland
--	Frank	Calgary
--	Jean	Montreal

CustomersInHouston

SSN	Name	City
234234	Mary	Houston

CustomersInSeattle

SSN	Name	City
345345	Sue	Seattle
345343	Joan	Seattle



.....

`CustomersInHouston(ssn,name,city)`
`CustomersInSeattle(ssn,name,city)`

`Customers(ssn,name,city)`

.....

Horizontal Partitioning

```
CREATE VIEW Customers AS
    CustomersInHouston
    UNION ALL      combines rows!
    CustomersInSeattle
    UNION ALL
    ...
```

`CustomersInHouston(ssn,name,city)`
`CustomersInSeattle(ssn,name,city)`

`Customers(ssn,name,city)`

.....

Horizontal Partitioning

```
SELECT name  
FROM   Customers  
WHERE  city = 'Seattle'
```

Which tables are inspected by the system ?

`CustomersInHouston(ssn,name,city)`
`CustomersInSeattle(ssn,name,city)`

`Customers(ssn,name,city)`

.....

Horizontal Partitioning

Better: remove `CustomerInHouston.city` etc

with horizontal partitioning, the database has to look through ALL tuples still because we do not split by attributes; BUT we can manually specify these qualities of the tables (i.e. that `CustomersInHouston` only have city = Houston)

```
CREATE VIEW Customers AS
  (SELECT SSN, name, 'Houston' as city
   FROM CustomersInHouston)
    UNION ALL
  (SELECT SSN, name, 'Seattle' as city
   FROM CustomersInSeattle)
    UNION ALL
  ...
```

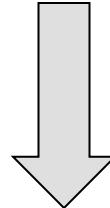
`CustomersInHouston(ssn,name,city)`
`CustomersInSeattle(ssn,name,city)`

`Customers(ssn,name,city)`

.....

Horizontal Partitioning

```
SELECT name  
FROM Customers  
WHERE city = 'Seattle'
```



```
SELECT name  
FROM CustomersInSeattle
```

Horizontal Partitioning Applications

- Performance optimization
 - Especially for data warehousing
 - E.g., one partition per month
 - E.g., archived applications and active applications
- Distributed and parallel databases
- Data integration

Conclusion

- Poor schemas can lead to performance inefficiencies
- E/R diagrams are means to structurally visualize and design relational schemas
- Normalization is a principled way of converting schemas into a form that avoid such problems
- BCNF is one of the most widely used normalized form in practice