

Introduction to Database Systems CSE 414

Lecture 27: Implementation of Transactions

CSE 414 - Autumn 2018

1

Announcements

- Final exam
 - Will test concepts from entire class but emphasis on post-midterm
 - Previous finals are for reference only, better to study lecture and section materials

CSE 414 - Autumn 2018

2

conflicts between a series of transactions

Testing for Conflict-Serializability

Precedence graph: which transactions have to go before others for process to be consistent

- A node for each transaction T_i ,
- An edge from T_i to T_j whenever an action in T_i conflicts with, and comes before an action in T_j

- The schedule is conflict-serializable iff the precedence graph is acyclic

if cyclic there is no possible order;
i.e. A needs to be run before B, B needs to be run before A = cycle = not possible

CSE 414 - Autumn 2018

3

Example 1

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$




CSE 414 - Autumn 2018

4

Example 1

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

read to A, the



This schedule is **conflict-serializable**

CSE 414 - Autumn 2018

5

Example 2

$r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B)$




CSE 414 - Autumn 2018

6

start with far left element $r_2(A)$ in this case:
There exists a conflict if we have a read write, write read, or write write conflict on SAME dataset
for each element:
Compare this element with all elements to right of it; if we have current element & element to the right of it
 $r_i(A) \& w_j(A)$
OR
 $w_i(A) \& r_j(A)$
OR
 $w_i(A) \& w_j(A)$

with $i \neq j$ then we have a conflict between i and j ; draw graph and check for cycles

Example 2



This schedule is **NOT** conflict-serializable

CSE 414 - Autumn 2018

7

Implementing Transactions

CSE 414 - Autumn 2018

8

preventing the formation of non-serializable conflicts (i.e. we can always find a proper order to run our transactions)

Scheduler

- **Scheduler** = the module that schedules the transaction's actions, ensuring serializability
- Also called **Concurrency Control Manager**
- We discuss next how a scheduler may be implemented

CSE 414 - Autumn 2018

9

Implementing a Scheduler

Major differences between database vendors

- **Locking Scheduler** starts from assumption that all transactions are interfering with each other.
 - Aka "pessimistic concurrency control"
 - SQLite, SQL Server, DB2
- **Multiversion Concurrency Control (MVCC)**
 - Aka "optimistic concurrency control"

We discuss only locking schedulers in this class

CSE 414 - Autumn 2018

10

Locking Scheduler

once a transaction acquires a lock, NO other transaction can access that element

Simple idea:

- Each element has a unique **lock**
- Each transaction must first **acquire** the lock before reading/writing that element
- If the lock is taken by another transaction, then wait
- The transaction must **release** the lock(s)

By using locks scheduler ensures conflict-serializability

prevents conflicts from occurring in the first place?

What Data Elements are Locked?

Major differences between vendors:

- Lock on the entire database
 - SQLite only one transaction can run on entire database at a time.
- Lock on individual records
 - SQL Server, DB2, etc

CSE 414 - Autumn 2018

12

More Notations

$L_i(A)$ = transaction T_i acquires lock for element A

$U_i(A)$ = transaction T_i releases lock for element A

CSE 414 - Autumn 2018

13

A Non-Serializable Schedule

T1	T2
READ(A)	
A := A+100	
WRITE(A)	
	READ(A)
	A := A*2
	WRITE(A)
	READ(B)
	B := B*2
	WRITE(B)
READ(B)	
B := B+100	
WRITE(B)	

CSE 414 - Autumn 2018

14

Example

T1	T2
$L_1(A)$; READ(A)	
A := A+100	
WRITE(A); $U_1(A)$; $L_1(B)$	
	$L_2(A)$; READ(A)
	A := A*2
	WRITE(A); $U_2(A)$;
	$L_2(B)$; BLOCKED...
READ(B)	
B := B+100	
WRITE(B); $U_1(B)$;	
	...GRANTED ; READ(B)
	B := B*2
	WRITE(B); $U_2(B)$;

Scheduler has ensured a conflict-serializable schedule

15

But what if...

T1	T2
$L_1(A)$; READ(A)	
A := A+100	
WRITE(A); $U_1(A)$;	
	$L_2(A)$; READ(A)
	A := A*2
	WRITE(A); $U_2(A)$;
	$L_2(B)$; READ(B)
	B := B*2
	WRITE(B); $U_2(B)$;
$L_1(B)$; READ(B)	
B := B+100	
WRITE(B); $U_1(B)$;	

Locks did not enforce conflict-serializability !!! What's wrong ?

16

Two Phase Locking (2PL)

The 2PL rule:

In every transaction, all lock requests must precede all unlock requests

CSE 414 - Autumn 2018

17

Example: 2PL transactions

T1	T2
$L_1(A)$; $L_1(B)$; READ(A)	
A := A+100	
WRITE(A); $U_1(A)$	
	$L_2(A)$; READ(A)
	A := A*2
	WRITE(A);
	$L_2(B)$; BLOCKED...
READ(B)	
B := B+100	
WRITE(B); $U_1(B)$;	
	...GRANTED ; READ(B)
	B := B*2
	WRITE(B); $U_2(A)$; $U_2(B)$;

Now it is conflict-serializable

CSE 414 - Autumn 2018

18

Two Phase Locking (2PL)


Theorem: 2PL ensures conflict serializability

19

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability


Proof. Suppose not: then there exists a cycle in the precedence graph.



Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.




Then there is the following **temporal** cycle in the schedule:

21

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.



Then there is the following **temporal** cycle in the schedule:
 $U_1(A) \rightarrow L_2(A)$ why?

transaction one has to unlock its element before any other ones can start on A


$U_1(A)$ happened strictly **before** $L_2(A)$

22

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.



Then there is the following **temporal** cycle in the schedule:
 $U_1(A) \rightarrow L_2(A)$
 $L_2(A) \rightarrow U_2(B)$ why?

by 2PL all locks occur BEFORE all unlocks


$L_2(A)$ happened strictly **before** $U_1(A)$

23

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.




Then there is the following **temporal** cycle in the schedule:
 $U_1(A) \rightarrow L_2(A)$
 $L_2(A) \rightarrow U_2(B)$ why?

24

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.



Then there is the following **temporal** cycle in the schedule:


$U_1(A) \rightarrow L_2(A)$
 $L_2(A) \rightarrow U_2(B)$
 $U_2(B) \rightarrow L_3(B)$ why?

25

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.



Then there is the following **temporal** cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$
 $L_2(A) \rightarrow U_2(B)$
 $U_2(B) \rightarrow L_3(B)$

.....etc.....


26

transaction aborts; when this happens ALL elements have to be reset to exactly what they were before the transaction began... But this will overwrite the work/writes of other transactions.

Two Phase Locking (2PL)

Theorem: 2PL ensures conflict serializability

Proof. Suppose not: then there exists a cycle in the precedence graph.



Then there is the following **temporal** cycle in the schedule:

$U_1(A) \rightarrow L_2(A)$
 $L_2(A) \rightarrow U_2(B)$
 $U_2(B) \rightarrow L_3(B)$
 $L_3(B) \rightarrow U_3(C)$
 $U_3(C) \rightarrow L_1(C)$
 $L_1(C) \rightarrow U_1(A)$

Cycle in time:
Contradiction

Text

A New Problem: Non-recoverable Schedule

T1	T2
$L_1(A); L_1(B); \text{READ}(A)$	
$A := A + 100$	
$\text{WRITE}(A); U_1(A)$	
	$L_2(A); \text{READ}(A)$
	$A := A * 2$
	$\text{WRITE}(A);$
	$L_2(B); \text{BLOCKED}...$
$\text{READ}(B)$	
$B := B + 100$	
$\text{WRITE}(B); U_1(B);$	
	$... \text{GRANTED}; \text{READ}(B)$
	$B := B * 2$
	$\text{WRITE}(B); U_2(A); U_2(B);$
	Commit

Rollback

CSE 414 - Autumn 2018

28

A New Problem: Non-recoverable Schedule

T1	T2
$L_1(A); L_1(B); \text{READ}(A)$	
$A := A + 100$	
$\text{WRITE}(A); U_1(A)$	
	$L_2(A); \text{READ}(A)$
	$A := A * 2$
	$\text{WRITE}(A);$
	$L_2(B); \text{BLOCKED}...$
$\text{READ}(B)$	
$B := B + 100$	
$\text{WRITE}(B); U_1(B);$	
	$... \text{GRANTED}; \text{READ}(B)$
	$B := B * 2$
	$\text{WRITE}(B); U_2(A); U_2(B);$
	Commit

Rollback

Elements A, B written by T1 are restored to their original value.

Autumn 2018

29

A New Problem: Non-recoverable Schedule

T1	T2
$L_1(A); L_1(B); \text{READ}(A)$	
$A := A + 100$	
$\text{WRITE}(A); U_1(A)$	
	$L_2(A); \text{READ}(A)$
	$A := A * 2$
	$\text{WRITE}(A);$
	$L_2(B); \text{BLOCKED}...$
$\text{READ}(B)$	
$B := B + 100$	
$\text{WRITE}(B); U_1(B);$	
	$... \text{GRANTED}; \text{READ}(B)$
	$B := B * 2$
	$\text{WRITE}(B); U_2(A); U_2(B);$
	Commit

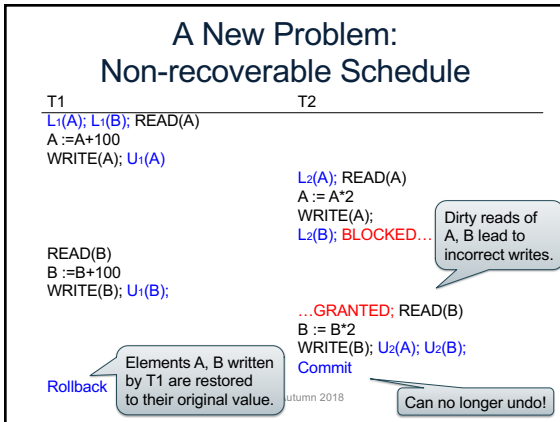
Rollback

Elements A, B written by T1 are restored to their original value.

Autumn 2018

30

Dirty reads of A, B lead to incorrect writes.



Strict 2PL

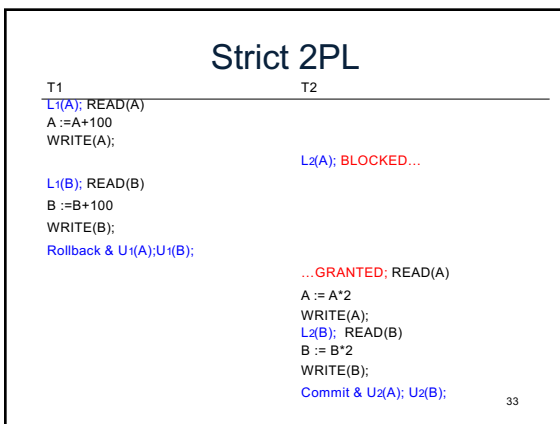
The Strict 2PL rule:

All locks are held until commit/abort:
 All unlocks are done together with commit/abort.

locks can occur anywhere throughout transaction, but unlocks **MUST** occur at end of transaction

With strict 2PL, we will get schedules that are both conflict-serializable and recoverable

CSE 414 - Autumn 2018 32



- ## Strict 2PL
- Lock-based systems always use strict 2PL
 - Easy to implement:
 - Before a transaction reads or writes an element A, insert an L(A) DONT just lock all elements at start of transaction
 - When the transaction commits/aborts, then release all locks
 - Ensures both conflict serializability and recoverability
- CSE 414 - Autumn 2018 34

Another problem: Deadlocks

- T₁: R(A), W(B)
- T₂: R(B), W(A)

- T₁ holds the lock on A, waits for B
- T₂ holds the lock on B, waits for A

one solution: monitor how long transactions last; abort slow running transactions

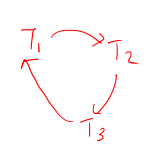
This is a deadlock!

CSE 414 - Autumn 2018 35

Another problem: Deadlocks

To detect a deadlocks, search for a cycle in the *waits-for graph*:

- T₁ waits for a lock held by T₂;
- T₂ waits for a lock held by T₃;
- ...
- T_n waits for a lock held by T₁



Relatively expensive: check periodically, if deadlock is found, then abort one transaction.
 need to continuously re-check for deadlocks

CSE 414 - Autumn 2018 36

A "Solution": Lock Modes

two different types of locks

- S = shared lock (for READ)
- X = exclusive lock (for WRITE)

Lock compatibility matrix:

	None	S	X
None			
S			
X			

CSE 414 - Autumn 2018

37

A "Solution": Lock Modes

- S = shared lock (for READ)
- X = exclusive lock (for WRITE)

Lock compatibility matrix:

	None	S	X
None			
S			
X			

CSE 414 - Autumn 2018

38

Lock Granularity

what do you lock? How big/small? DB or single tuples?

lock individual tuples/attributes; more likely to create deadlocks

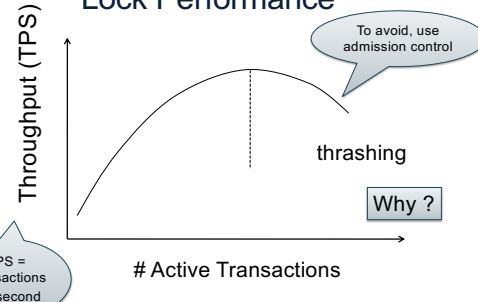
- **Fine granularity locking** (e.g., tuples)
 - High concurrency
 - High overhead in managing locks
 - E.g., SQL Server
- lock entire tables or entire database; prevents deadlocks but slow, as you lock a lot of data you really didn't need to
- **Coarse grain locking** (e.g., tables, entire database)
 - Many false conflicts
 - Less overhead in managing locks
 - E.g., SQL Lite

- Solution: lock escalation changes granularity as needed

CSE 414 - Autumn 2018

39

Lock Performance



CSE 414 - Autumn 2018

40

Phantom Problem

- So far we have assumed the database to be a *static* collection of elements (=tuples)
- If tuples are inserted/deleted then the *phantom problem* appears

CSE 414 - Autumn 2018

41

Suppose there are two blue products, A1, A2:

Phantom Problem

<p>T1</p> <pre>SELECT * FROM Product WHERE color='blue'</pre> <p>SELECT *</p> <pre>FROM Product WHERE color='blue'</pre>	<p>T2</p> <pre>INSERT INTO Product(name, color) VALUES ('A3','blue')</pre>
--	--

Is this schedule serializable ?

CSE 414 - Autumn 2018

42

Suppose there are two blue products, A1, A2:

Phantom Problem

T1	T2
SELECT *	
FROM Product	
WHERE color='blue'	
	INSERT INTO Product(name, color)
	VALUES ('A3','blue')
SELECT *	
FROM Product	
WHERE color='blue'	

R₁(A1);R₁(A2);W₂(A3);R₁(A1);R₁(A2);R₁(A3)

CSE 414 - Autumn 2018

43

Suppose there are two blue products, A1, A2:

Phantom Problem

T1	T2
SELECT *	
FROM Product	
WHERE color='blue'	
	INSERT INTO Product(name, color)
	VALUES ('A3','blue')
SELECT *	
FROM Product	
WHERE color='blue'	

R₁(A1);R₁(A2);W₂(A3);R₁(A1);R₁(A2);R₁(A3)

W₂(A3);R₁(A1);R₁(A2);R₁(A1);R₁(A2);R₁(A3)⁴⁴

CSE 414 - Autumn 2018

Phantom Problem

- A "phantom" is a tuple that is invisible during **part** of a transaction execution but not invisible during the **entire** execution
- In our example:
 - T1: reads list of products
 - T2: inserts a new product
 - T1: re-reads: a new product appears !



CSE 414 - Autumn 2018

45

Dealing With Phantoms

- Lock the entire table
- Lock the index entry for 'blue'
 - If index is available
- Or use predicate locks
 - A lock on an arbitrary predicate

Dealing with phantoms is expensive !

CSE 414 - Autumn 2018

46

Summary of Serializability

- Serializable schedule = equivalent to a serial schedule
- (strict) 2PL guarantees *conflict serializability*
 - What is the difference?
- Static database:**
 - Conflict serializability* implies serializability
- Dynamic database:**
 - This no longer holds

CSE 414 - Autumn 2018

47

Isolation Levels in SQL

- "Dirty reads"

SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED
- "Committed reads"

SET TRANSACTION ISOLATION LEVEL READ COMMITTED
- "Repeatable reads"

SET TRANSACTION ISOLATION LEVEL REPEATABLE READ
- Serializable transactions

SET TRANSACTION ISOLATION LEVEL SERIALIZABLE

ACID

CSE 414 - Autumn 2018

48

1. Isolation Level: Dirty Reads

- "Long duration" WRITE locks
 - Strict 2PL
- No READ locks
 - Read-only transactions are never delayed

Possible problems: dirty and inconsistent reads

CSE 414 - Autumn 2018

49

2. Isolation Level: Read Committed

- "Long duration" WRITE locks
 - Strict 2PL
- "Short duration" READ locks
 - Only acquire lock while reading (not 2PL)

Unrepeatable reads:
When reading same element twice,
may get two different values

CSE 414 - Autumn 2018

50

3. Isolation Level: Repeatable Read

- "Long duration" WRITE locks
 - Strict 2PL
- "Long duration" READ locks
 - Strict 2PL

Why ?

This is not serializable yet !!!

CSE 414 - Autumn 2018

51

4. Isolation Level Serializable

- "Long duration" WRITE locks
 - Strict 2PL
- "Long duration" READ locks
 - Strict 2PL
- Predicate locking
 - To deal with phantoms

CSE 414 - Autumn 2018

52

Beware!

In commercial DBMSs:

- Default level is often NOT serializable
- Default level differs between DBMSs
- Some engines support subset of levels!
- Serializable may not be exactly ACID
 - Locking ensures isolation, not atomicity
- Also, some DBMSs do NOT use locking and different isolation levels can lead to different pbs
- Bottom line: RTFM for your DBMS!

CSE 414 - Autumn 2018

53